

# Homework 3: Convolutional Neural Networks

Due Wednesday 11/24 at 11:59 pm EST

Download the dataset `cats-notcats` from github (given as a part of the assignment). This dataset has images of cats and images that are not cats (in separate folders). The task is to train a convolutional neural network (CNN) to build a classifier that can classify a new image as either `cat` or `not cat`

1. Load the dataset and create three stratified splits - train/validation/test in the ratio of 70/10/20.

```
In [4]: import tensorflow as tf
        from tensorflow import keras

        batch_size=32
        image_generator = tf.keras.preprocessing.image.ImageDataGenerator(validation
all_data_batches = image_generator.flow_from_directory(directory='/Users/Gr
                                                    class_mode='binary',

        val_gen = image_generator.flow_from_directory(directory='/Users/Griffin/rep
                                                    class_mode='binary', target_s

        #print(batches)
```

```
Found 5668 images belonging to 2 classes.
Found 1133 images belonging to 2 classes.
```

```
In [5]: #for batch in all_data_batches:
        #     print(batch)
```

2. Create a CNN that has the following hidden layers:
  - a. 2D convolution layer with a 3x3 kernel size, has 128 filters, stride of 1 and padded to yield the same size as input, followed by a ReLU activation layer
  - b. Max pooling layer of 2x2
  - c. Dense layer with 128 dimensions and ReLU as the activation layer

```
In [6]: model = tf.keras.Sequential(  
    [  
        tf.keras.layers.Conv2D(128, kernel_size=(3,3), activation='relu',  
        tf.keras.layers.MaxPooling2D(  
            pool_size=(2,2)),  
        tf.keras.layers.Flatten(),  
        tf.keras.layers.Dense(128, activation="relu", name="c"),  
    ]  
)  
model.compile(loss='binary_crossentropy',  
              optimizer='rmsprop',  
              metrics=[ 'accuracy' ])
```

3. Train the classifier for 20 epochs with 100 steps per epoch. Also use the validation data during training the estimator.

```
In [7]: model.fit_generator(all_data_batches, epochs=20, steps_per_epoch=100, validation_data=val_gen, model.save_weights(first_try.h5))
```

```
<ipython-input-7-13f639a6401b>:1: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.
```

```
model.fit_generator(all_data_batches, epochs=20, steps_per_epoch=100, validation_data=val_gen)
```

```
-----
--
ValueError                                Traceback (most recent call last)
```

```
<ipython-input-7-13f639a6401b> in <module>
```

```
----> 1 model.fit_generator(all_data_batches, epochs=20, steps_per_epoch=100, validation_data=val_gen)
```

```
      2 model.save_weights(first_try.h5)
```

```
~/opt/anaconda3/lib/python3.8/site-packages/keras/engine/training.py in fit_generator(self, generator, steps_per_epoch, epochs, verbose, callbacks, validation_data, validation_steps, validation_freq, class_weight, max_queue_size, workers, use_multiprocessing, shuffle, initial_epoch)
```

```
    2014         'Please use `Model.fit`, which supports generators.',
```

```
    2015         stacklevel=2)
```

```
-> 2016     return self.fit(
```

```
    2017         generator,
```

```
    2018         steps_per_epoch=steps_per_epoch,
```

```
~/opt/anaconda3/lib/python3.8/site-packages/keras/utils/traceback_utils.py in error_handler(*args, **kwargs)
```

```
    65     except Exception as e: # pylint: disable=broad-except
```

```
    66         filtered_tb = _process_traceback_frames(e.__traceback__)
```

```
----> 67         raise e.with_traceback(filtered_tb) from None
```

```
    68     finally:
```

```
    69         del filtered_tb
```

```
~/opt/anaconda3/lib/python3.8/site-packages/keras/engine/input_spec.py in assert_input_compatibility(input_spec, inputs, layer_name)
```

```
    225     ndim = x.shape.rank
```

```
    226     if ndim is not None and ndim < spec.min_ndim:
```

```
--> 227         raise ValueError(f'Input {input_index} of layer "{layer_name}" ' 'is incompatible with the layer: ' 'expected min_ndim={spec.min_ndim}, ' 'found ndim={ndim}. Full shape received: {x.shape}')
```

```
    228
```

```
    229         f'expected min_ndim={spec.min_ndim}, ' 'found ndim={ndim}. Full shape received: {x.shape}')
```

```
ValueError: Exception encountered when calling layer "sequential" (type Sequential).

Input 0 of layer "conv2d" is incompatible with the layer: expected min_ndim=4, found ndim=2. Full shape received: (16, 196608)
```

```
Call arguments received:
```

- inputs=tf.Tensor(shape=(16, 256, 256, 3), dtype=float32)
- training=False
- mask=None

## 4. Plot the accuracy and the loss over epochs for train &amp; validation sets

In [ ]:

## 5. Add the following layers to (2) before the dense layer:

- 2D convolution layer with a 3x3 kernel size, has 64 filters, stride of 1 and padded to yield the same size as input, followed by a ReLU activation layer
- Max pooling layer of 2x2
- 2D convolution layer with a 3x3 kernel size, has 32 filters, stride of 1 and padded to yield the same size as input, followed by a ReLU activation layer
- Max pooling layer of 2x2
- Dense layer with 256 dimensions and ReLU as the activation layer

In [17]:

```
model2 = tf.keras.Sequential(
    [
        tf.keras.layers.Conv2D(kernel_size=(3,3), activation='relu', padding='same',
                                strides=1, input_shape=(256,256,3)),
        tf.keras.layers.MaxPooling2D(
            pool_size=(2,2)),
        tf.keras.layers.Conv2D(kernel_size=(3,3), activation='relu', padding='same',
                                strides=1, input_shape=(256,256,3)),
        tf.keras.layers.MaxPooling2D(
            pool_size=(2,2)),
        tf.keras.layers.Dense(256, activation="relu"),
        tf.keras.layers.Dense(128, activation="relu"),
    ]
)
model2.compile(tf.keras.optimizers.Adam(lr=0.01), loss="categorical_crossentropy")
```

## 6. Train the classifier again for 20 epochs with 100 steps per epoch. Also use the validation data during training the estimator.

In [19]: `model.fit(INPUT, OUTPUT, epochs=20, steps_per_epoch=100)`

## 7. Plot the accuracy and the loss over epochs for train &amp; validation sets

In [21]: `#code here`