

W4995 Applied Machine Learning

Fall 2021

Lecture 9
Dr. Vijay Pappu

Announcements

- HW2 grades are released
- Project deliverable 2 grades are released
- Project deliverable 3 does not require a presentation and is due on 12/15
- Midterms are being graded currently
- No class next week (Thanksgiving break)

In today's lecture, we will cover...

- Convolutional Neural Networks
- Recurrent Neural Networks

Convolutional Neural Networks

Convolutional Neural Networks

- Convolutional neural networks (CNNs) have been highly successful in many practical applications
 - Image classification
 - Object detection
 - ...
- The term “Convolutional Neural Network” indicates that the network employs a mathematical operation called “**convolution**”
- “Convolution” is a specialized kind of linear operator
- CNNs are simply neural networks that use convolution instead of general matrix multiplication in at least one of their layers.

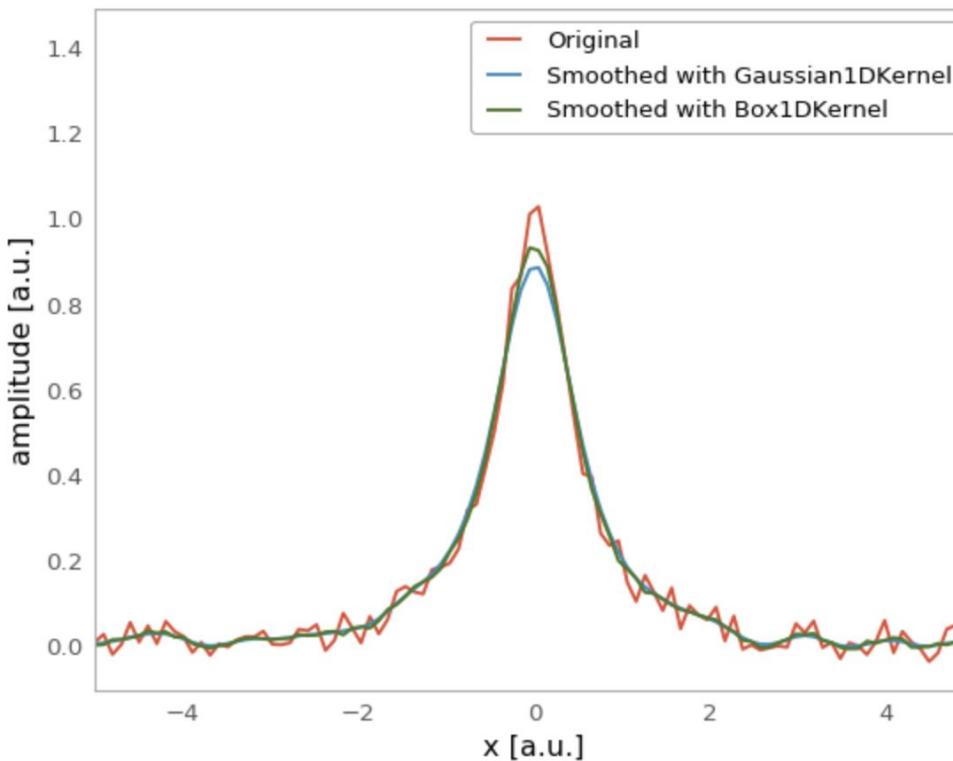
Convolution - 1D function

$$s(t) = \int x(a)w(t-a)da.$$

$$s(t) = (x * w)(t).$$

$$s(t) = (x * w)(t) = \sum_{a=-\infty}^{\infty} x(a)w(t-a).$$

Convolution - 1D function



Convolution - 2D function

$$S(i, j) = (I * K)(i, j) = \sum_m \sum_n I(m, n)K(i - m, j - n).$$

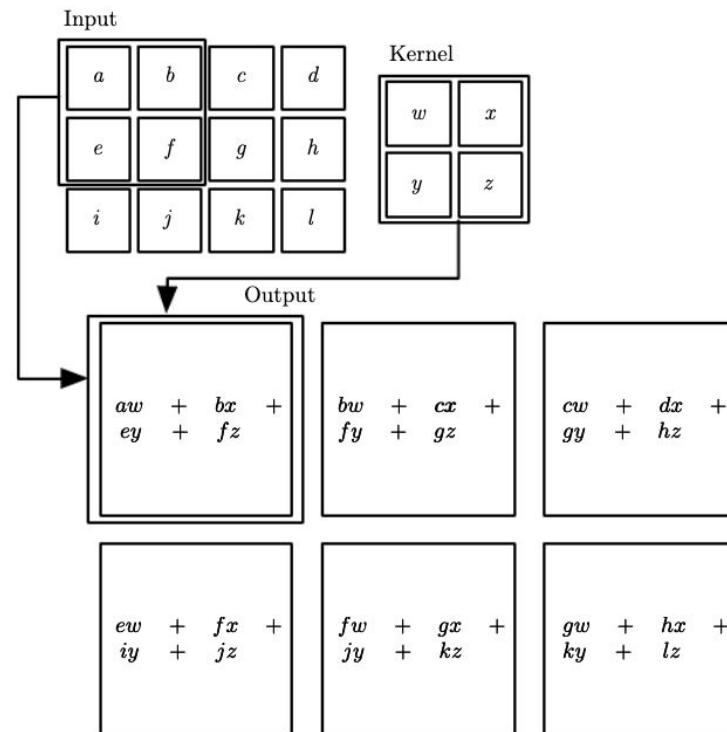
$$S(i, j) = (K * I)(i, j) = \sum_m \sum_n I(i - m, j - n)K(m, n).$$

$$S(i, j) = (K * I)(i, j) = \sum_m \sum_n I(i + m, j + n)K(m, n).$$

cross-correlation

(2D convolution without kernel flipping)

Convolution - 2D function



Convolution - 2D function

| | | | | |
|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Input

| | | |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |

Filter / Kernel

Convolution - 2D function

| | | | | |
|-----|-----|-----|---|---|
| 1x1 | 1x0 | 1x1 | 0 | 0 |
| 0x0 | 1x1 | 1x0 | 1 | 0 |
| 0x1 | 0x0 | 1x1 | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Input

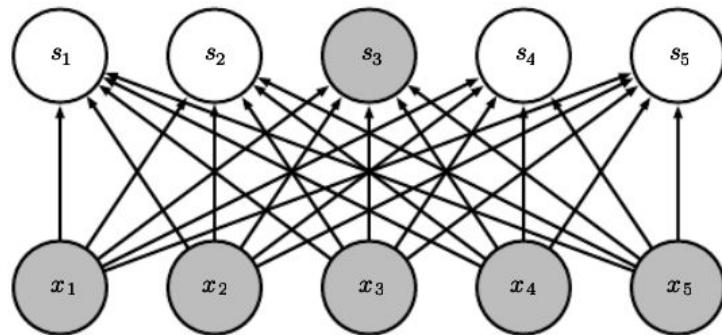
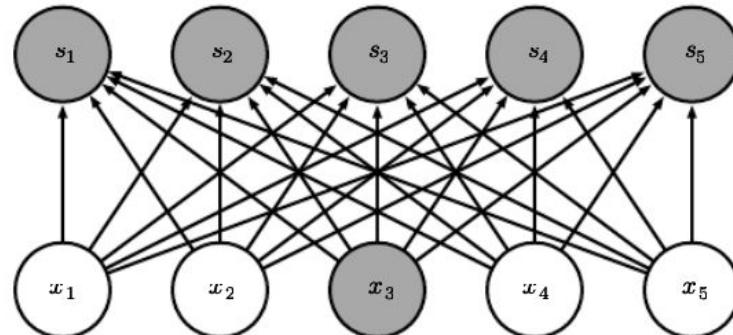
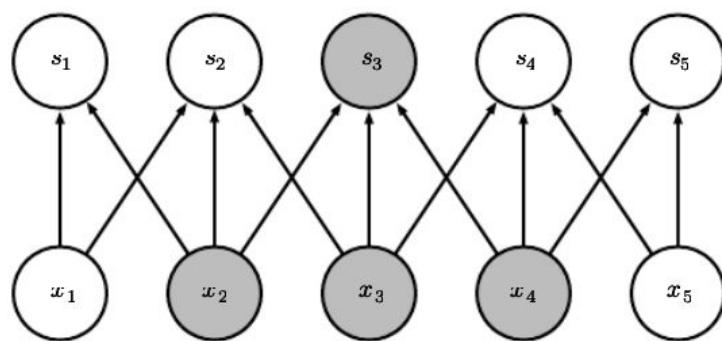
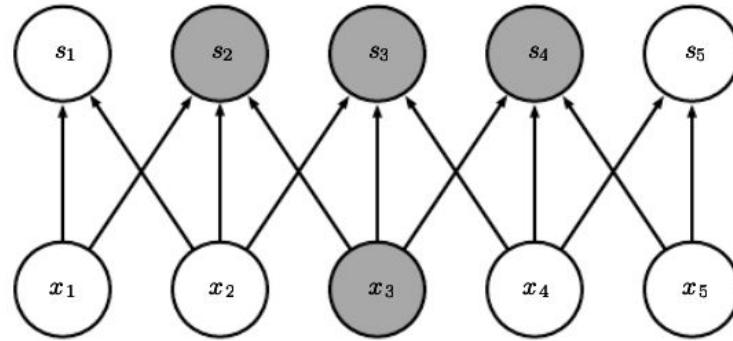
| | | |
|---|--|--|
| 4 | | |
| | | |
| | | |

Feature map

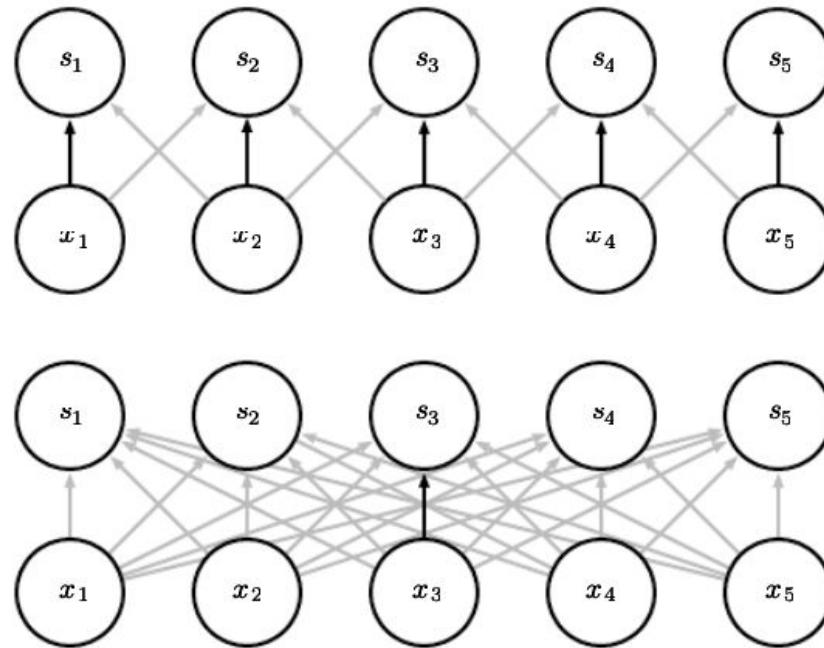
Convolutions in Neural Networks

- Sparse interactions
- Parameter sharing
- Equivariant representations

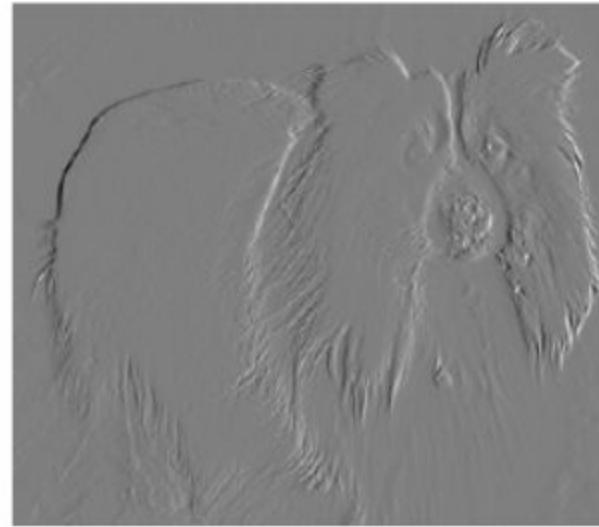
Convolutions in Neural Networks - Sparse Interactions



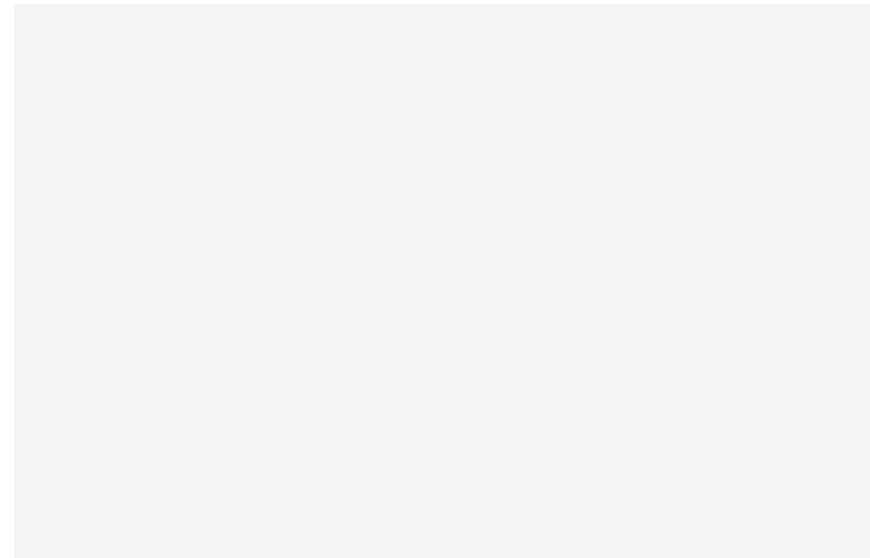
Convolutions in Neural Networks - Parameter Sharing



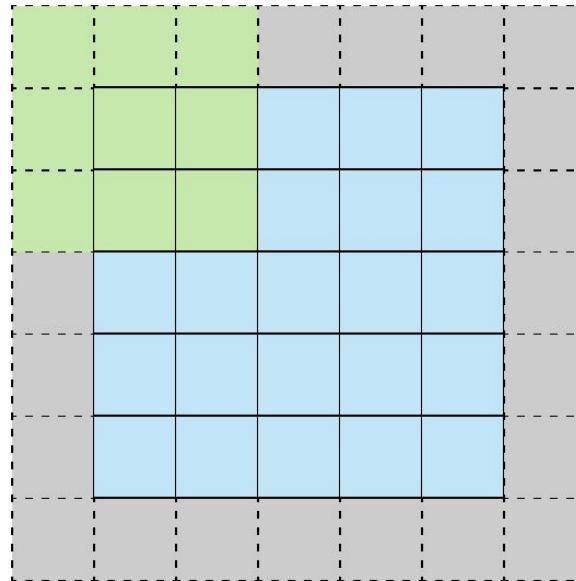
Convolutions in Neural Networks



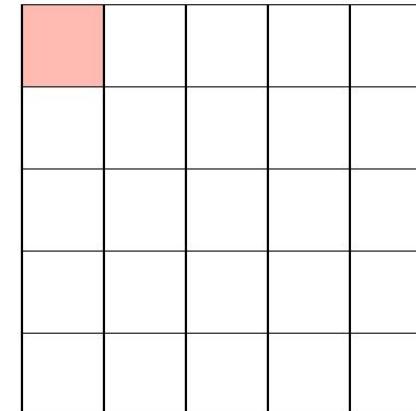
Convolutions in Neural Networks - Stride & Padding



Convolutions in Neural Networks - Stride & Padding



Stride 1 with Padding



Feature Map

Convolutions in Neural Networks - Activation function

| | | | | |
|---|---|-----|-----|-----|
| 1 | 1 | 1 | 0 | 0 |
| 0 | 1 | 1 | 1 | 0 |
| 0 | 0 | 1x1 | 1x0 | 1x1 |
| 0 | 0 | 1x0 | 1x1 | 0x0 |
| 0 | 1 | 1x1 | 0x0 | 0x1 |

| | | |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| 2 | 3 | 4 |

→ $h(x)$

Convolutional Neural Networks - Pooling

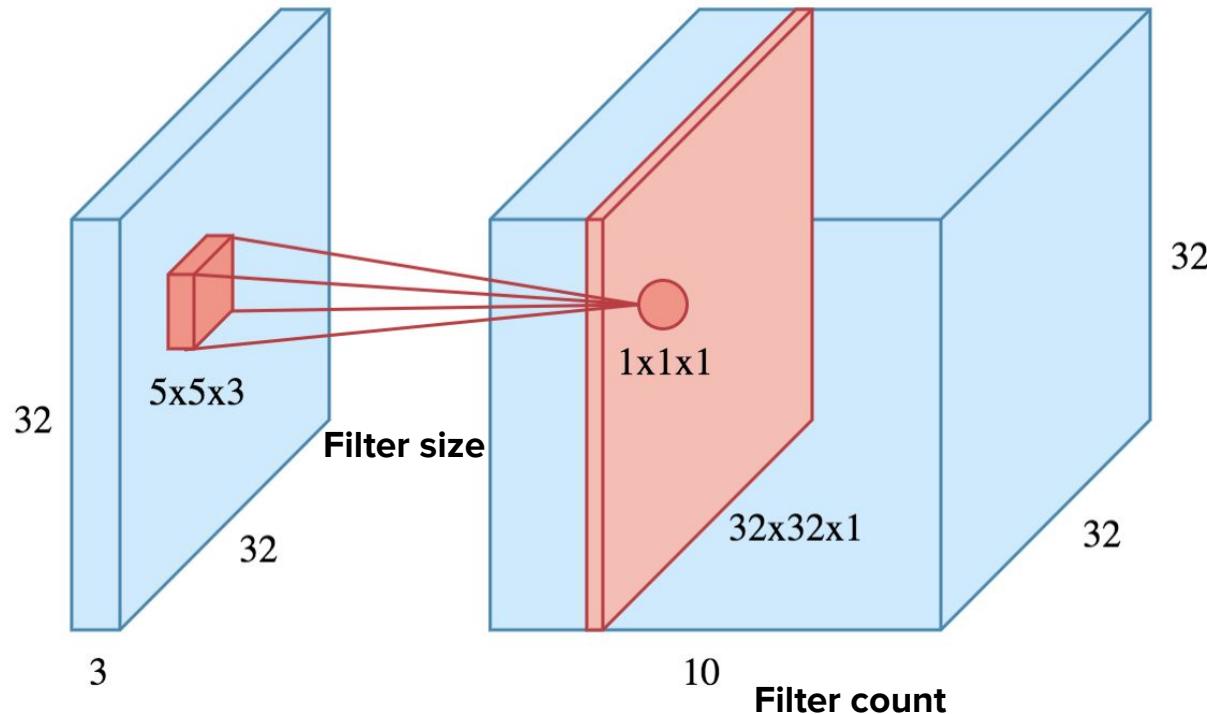
| | | | |
|---|---|---|---|
| 1 | 1 | 2 | 4 |
| 5 | 6 | 7 | 8 |
| 3 | 2 | 1 | 0 |
| 1 | 2 | 3 | 4 |

max pool with 2x2
window and stride 2

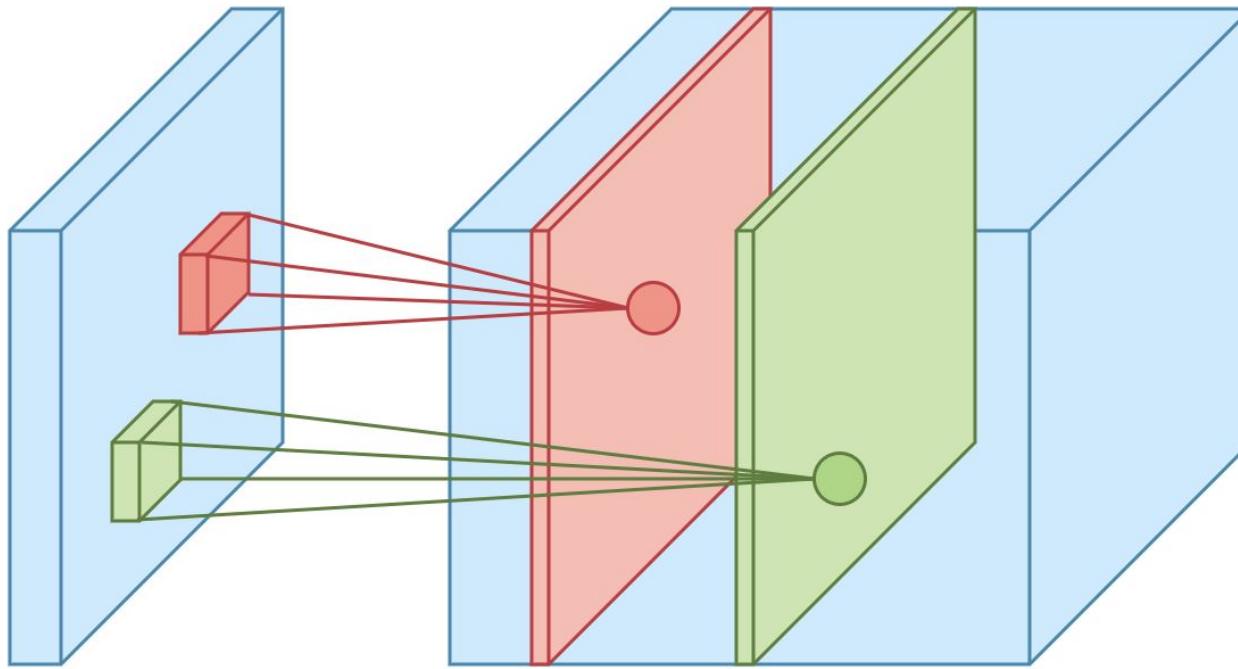
The diagram illustrates a 2x2 max pooling operation with a stride of 2. The input matrix is a 4x4 grid of numbers. A 2x2 window slides across the input with a stride of 2, and the maximum value in each window is selected as the value in the corresponding position of the output matrix. The output matrix has a 2x2 grid of values: 6 (top-left), 8 (top-right), 3 (bottom-left), and 4 (bottom-right).

| | |
|---|---|
| 6 | 8 |
| 3 | 4 |

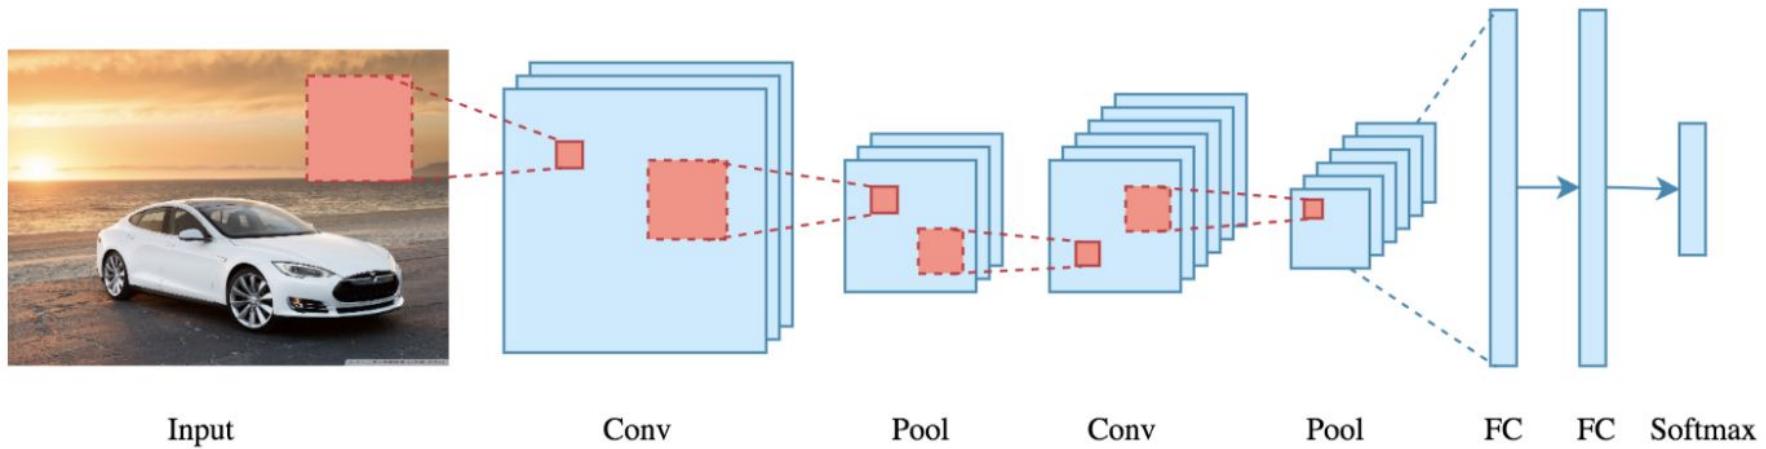
3D Convolutions in Neural Networks



3D Convolutions in Neural Networks



Typical CNN architecture



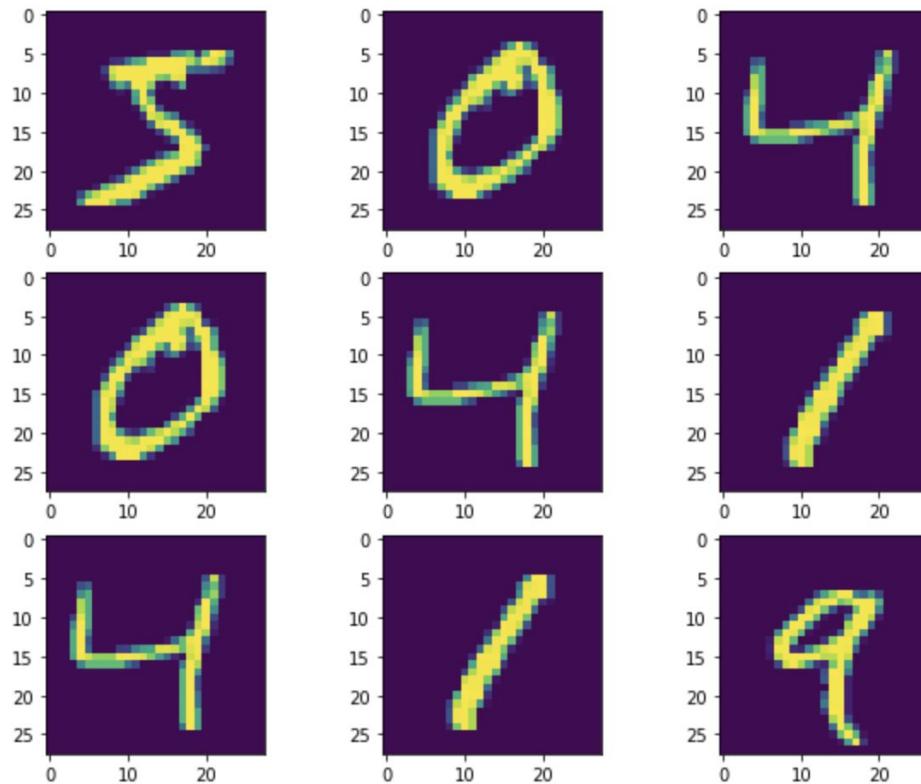
Training Convolutional Neural Networks

- CNNs are generally trained using backpropagation algorithm
- Math gets messy

Convolutional Neural Networks - Hyperparameters

- Network architecture
- Convolutional layer
 - Filter size
 - Filter count
 - Padding (mostly yes)
 - Stride (usually 1)
- Optimization algorithm parameters
- Activation function (ReLU, tanh, sigmoid etc.)
- Weight initialization
- Batch size
- # of epochs

Convolutional Neural Networks - Example



Convolutional Neural Networks - Example

```
batch_size = 128
num_classes = 10
# input image dimensions
img_rows, img_cols = 28, 28
# the data, shuffled and split between train and test sets
(x_dev, y_dev), (x_test, y_test) = mnist.load_data()
X_dev_images = x_dev.reshape(x_dev.shape[0], img_rows, img_cols, 1)
X_test_images = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
input_shape = (img_rows, img_cols, 1)
y_dev = np_utils.to_categorical(y_dev, num_classes)
y_test = np_utils.to_categorical(y_test, num_classes)
```

Convolutional Neural Networks - Example

```
num_classes = 10
cnn = Sequential()
cnn.add(Conv2D(32, kernel_size=(3, 3),
              activation='relu',
              input_shape=input_shape))
cnn.add(MaxPooling2D(pool_size=(2, 2)))
cnn.add(Conv2D(32, (3, 3), activation='relu'))
cnn.add(MaxPooling2D(pool_size=(2, 2)))
cnn.add(Flatten())
cnn.add(Dense(64, activation='relu'))
cnn.add(Dense(num_classes, activation='softmax'))
```

Convolutional Neural Networks - Example

Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|---------------------------------|--------------------|---------|
| <hr/> | | |
| conv2d (Conv2D) | (None, 26, 26, 32) | 320 |
| max_pooling2d (MaxPooling2D) | (None, 13, 13, 32) | 0 |
| conv2d_1 (Conv2D) | (None, 11, 11, 32) | 9248 |
| max_pooling2d_1 (MaxPooling 2D) | (None, 5, 5, 32) | 0 |
| flatten_1 (Flatten) | (None, 800) | 0 |
| dense (Dense) | (None, 64) | 51264 |
| dense_1 (Dense) | (None, 10) | 650 |
| <hr/> | | |

Total params: 61,482

Trainable params: 61,482

Non-trainable params: 0

Convolutional Neural Networks - Example

Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|---------------------------------|--------------------|---------|
| conv2d (Conv2D) | (None, 26, 26, 32) | 320 |
| max_pooling2d (MaxPooling2D) | (None, 13, 13, 32) | 0 |
| conv2d_1 (Conv2D) | (None, 11, 11, 32) | 9248 |
| max_pooling2d_1 (MaxPooling 2D) | (None, 5, 5, 32) | 0 |
| flatten_1 (Flatten) | (None, 800) | 0 |
| dense (Dense) | (None, 64) | 51264 |
| dense_1 (Dense) | (None, 10) | 650 |

$$(3*3*1+1)*32 = 320$$

Total params: 61,482
Trainable params: 61,482
Non-trainable params: 0

Convolutional Neural Networks - Example

Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|---------------------------------|--------------------|---------|
| conv2d (Conv2D) | (None, 26, 26, 32) | 320 |
| max_pooling2d (MaxPooling2D) | (None, 13, 13, 32) | 0 |
| conv2d_1 (Conv2D) | (None, 11, 11, 32) | 9248 |
| max_pooling2d_1 (MaxPooling 2D) | (None, 5, 5, 32) | 0 |
| flatten_1 (Flatten) | (None, 800) | 0 |
| dense (Dense) | (None, 64) | 51264 |
| dense_1 (Dense) | (None, 10) | 650 |

$$(3*3*1+1)*32 = 320$$

$$(3*3*32+1)*32 = 9248$$

Total params: 61,482
Trainable params: 61,482
Non-trainable params: 0

Convolutional Neural Networks - Example

Model: "sequential_1"

| Layer (type) | Output Shape | Param # | |
|---------------------------------|--------------------|---------|------------------------|
| conv2d (Conv2D) | (None, 26, 26, 32) | 320 | $(3*3*1+1)*32 = 320$ |
| max_pooling2d (MaxPooling2D) | (None, 13, 13, 32) | 0 | |
| conv2d_1 (Conv2D) | (None, 11, 11, 32) | 9248 | $(3*3*32+1)*32 = 9248$ |
| max_pooling2d_1 (MaxPooling 2D) | (None, 5, 5, 32) | 0 | |
| flatten_1 (Flatten) | (None, 800) | 0 | |
| dense (Dense) | (None, 64) | 51264 | $(800+1)*64 = 51264$ |
| dense_1 (Dense) | (None, 10) | 650 | |
| <hr/> | | | |
| Total params: 61,482 | | | |
| Trainable params: 61,482 | | | |
| Non-trainable params: 0 | | | |

Convolutional Neural Networks - Example

Model: "sequential_1"

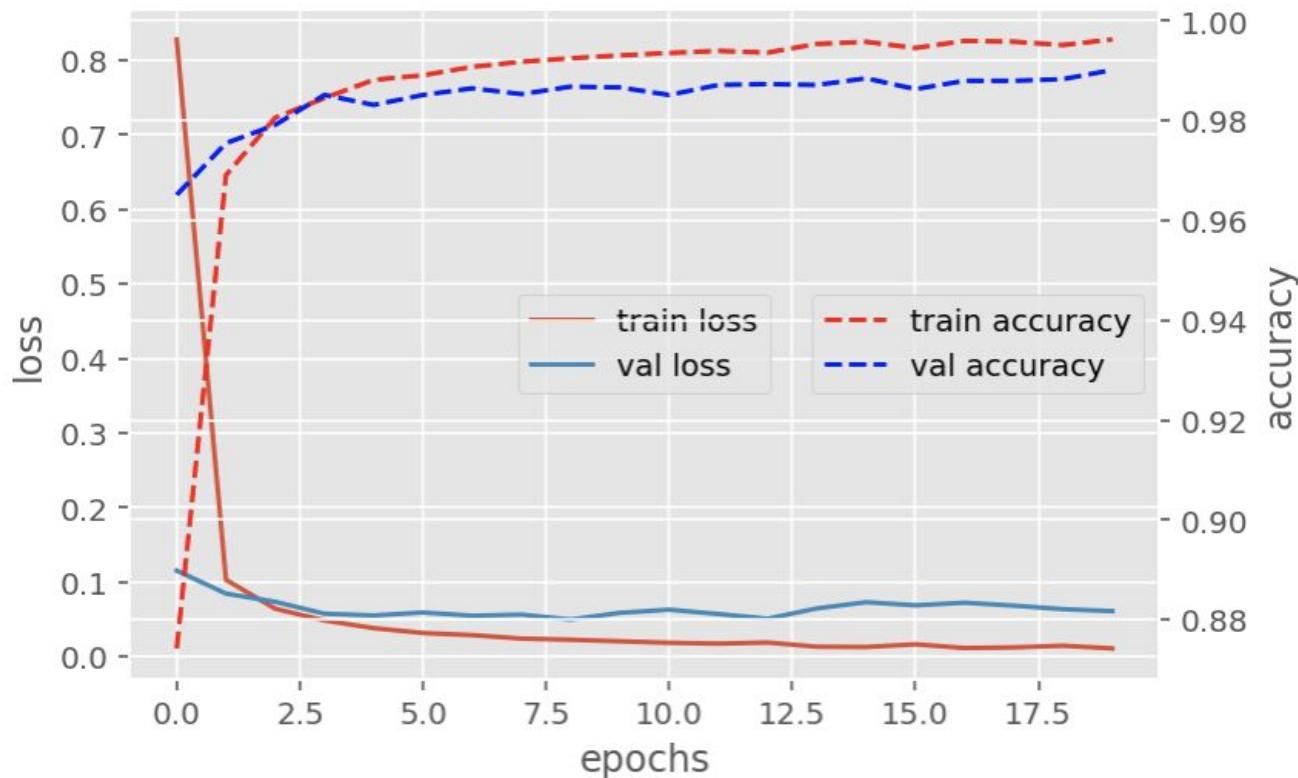
| Layer (type) | Output Shape | Param # | |
|---------------------------------|--------------------|---------|------------------------|
| conv2d (Conv2D) | (None, 26, 26, 32) | 320 | $(3*3*1+1)*32 = 320$ |
| max_pooling2d (MaxPooling2D) | (None, 13, 13, 32) | 0 | |
| conv2d_1 (Conv2D) | (None, 11, 11, 32) | 9248 | $(3*3*32+1)*32 = 9248$ |
| max_pooling2d_1 (MaxPooling 2D) | (None, 5, 5, 32) | 0 | |
| flatten_1 (Flatten) | (None, 800) | 0 | |
| dense (Dense) | (None, 64) | 51264 | $(800+1)*64 = 51264$ |
| dense_1 (Dense) | (None, 10) | 650 | $(64+1)*10 = 650$ |
| <hr/> | | | |
| Total params: 61,482 | | | |
| Trainable params: 61,482 | | | |
| Non-trainable params: 0 | | | |

Convolutional Neural Networks - Example

```
cnn.compile("adam", "categorical_crossentropy", metrics=['accuracy'])
history_cnn = cnn.fit(X_dev_images, y_dev,
                      batch_size=128, epochs=20, verbose=1, validation_split=.1)
cnn.evaluate(X_test_images, y_test)

Epoch 1/20
422/422 [=====] - 17s 38ms/step - loss: 0.8281 - accuracy: 0.8743 - val_loss: 0.1155 - val_accuracy: 0.
9652
Epoch 2/20
422/422 [=====] - 16s 37ms/step - loss: 0.1032 - accuracy: 0.9691 - val_loss: 0.0847 - val_accuracy: 0.
9755
Epoch 3/20
422/422 [=====] - 17s 40ms/step - loss: 0.0644 - accuracy: 0.9806 - val_loss: 0.0737 - val_accuracy: 0.
9792
Epoch 4/20
422/422 [=====] - 17s 40ms/step - loss: 0.0486 - accuracy: 0.9845 - val_loss: 0.0576 - val_accuracy: 0.
9852
Epoch 5/20
422/422 [=====] - 18s 42ms/step - loss: 0.0383 - accuracy: 0.9882 - val_loss: 0.0553 - val_accuracy: 0.
9832
Epoch 6/20
422/422 [=====] - 18s 43ms/step - loss: 0.0318 - accuracy: 0.9891 - val_loss: 0.0593 - val_accuracy: 0.
9852
Epoch 7/20
422/422 [=====] - 19s 46ms/step - loss: 0.0289 - accuracy: 0.9908 - val_loss: 0.0550 - val_accuracy: 0.
9865
Epoch 8/20
422/422 [=====] - 21s 49ms/step - loss: 0.0243 - accuracy: 0.9918 - val_loss: 0.0565 - val_accuracy: 0.
9853
```

Convolutional Neural Networks - Example



CNNs v.s. Fully Connected NNs - Example

```
model.summary()
```

Model: "sequential_3"

| Layer (type) | Output Shape | Param # |
|---------------------------|--------------|---------|
| dense_4 (Dense) | (None, 512) | 401920 |
| dense_5 (Dense) | (None, 10) | 5130 |
| <hr/> | | |
| Total params: 407,050 | | |
| Trainable params: 407,050 | | |
| Non-trainable params: 0 | | |

Fully connected NN

Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|--------------------------------|--------------------|---------|
| conv2d (Conv2D) | (None, 26, 26, 32) | 320 |
| max_pooling2d (MaxPooling2D) | (None, 13, 13, 32) | 0 |
| conv2d_1 (Conv2D) | (None, 11, 11, 32) | 9248 |
| max_pooling2d_1 (MaxPooling2D) | (None, 5, 5, 32) | 0 |
| flatten_1 (Flatten) | (None, 800) | 0 |
| dense (Dense) | (None, 64) | 51264 |
| dense_1 (Dense) | (None, 10) | 650 |
| <hr/> | | |

Total params: 61,482
Trainable params: 61,482
Non-trainable params: 0

Convolutional NN

Convolutional Neural Networks - Another Example



Convolutional Neural Networks - Another Example

```
base_dir = 'dogs-vs-cats-small'
train_dir = os.path.join(base_dir, 'train')
validation_dir = os.path.join(base_dir, 'validate')
test_dir = os.path.join(base_dir, 'test')

train_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(150, 150),
    batch_size=20,
    class_mode='binary')

validation_generator = test_datagen.flow_from_directory(
    validation_dir,
    target_size=(150, 150),
    batch_size=20,
    class_mode='binary')
```

Found 2000 images belonging to 2 classes.
Found 1000 images belonging to 2 classes.

Convolutional Neural Networks - Another Example

```
model = Sequential()
model.add(Conv2D(32, (3, 3), activation='relu', padding='same', name='conv_1',
                input_shape=(150, 150, 3)))
model.add(MaxPooling2D((2, 2), name='maxpool_1'))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', name='conv_2'))
model.add(MaxPooling2D((2, 2), name='maxpool_2'))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', name='conv_3'))
model.add(MaxPooling2D((2, 2), name='maxpool_3'))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', name='conv_4'))
model.add(MaxPooling2D((2, 2), name='maxpool_4'))
model.add(Flatten())
model.add(Dropout(0.5))
model.add(Dense(512, activation='relu', name='dense_1'))
model.add(Dense(256, activation='relu', name='dense_2'))
model.add(Dense(1, activation='sigmoid', name='output'))

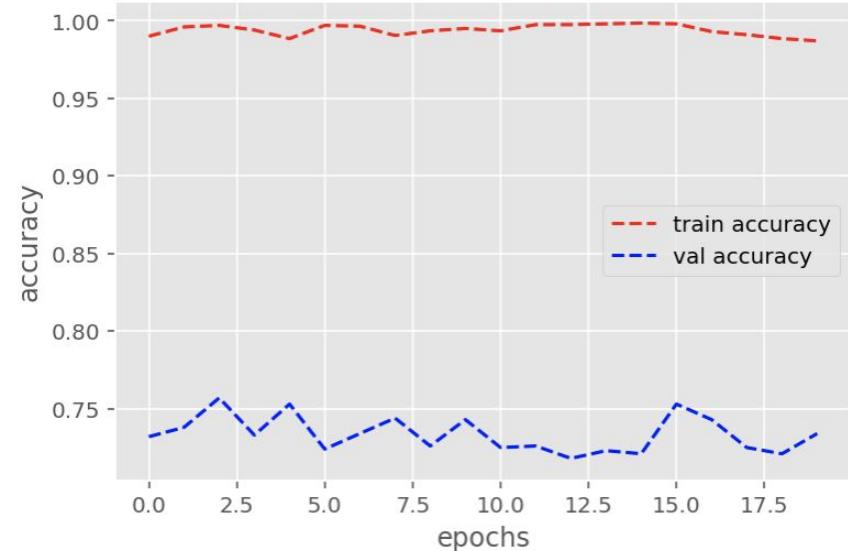
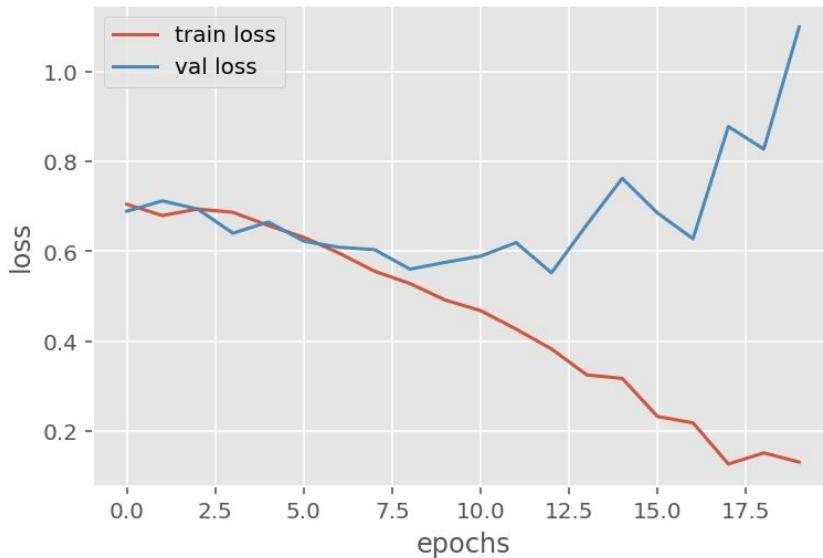
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

Convolutional Neural Networks - Another Example

```
history = model.fit_generator(train_generator, steps_per_epoch=100, epochs=20,
                               validation_data=validation_generator, validation_steps=50, verbose=1)

Epoch 1/20
100/100 [=====] - 27s 266ms/step - loss: 0.7045 - accuracy: 0.5100 - val_loss: 0.6891 - val_accuracy: 0.5800
Epoch 2/20
100/100 [=====] - 28s 275ms/step - loss: 0.6795 - accuracy: 0.5510 - val_loss: 0.7122 - val_accuracy: 0.5060
Epoch 3/20
100/100 [=====] - 29s 292ms/step - loss: 0.6939 - accuracy: 0.5335 - val_loss: 0.6939 - val_accuracy: 0.4790
Epoch 4/20
100/100 [=====] - 30s 297ms/step - loss: 0.6868 - accuracy: 0.5590 - val_loss: 0.6401 - val_accuracy: 0.6160
Epoch 5/20
100/100 [=====] - 28s 283ms/step - loss: 0.6573 - accuracy: 0.5875 - val_loss: 0.6652 - val_accuracy: 0.6110
Epoch 6/20
100/100 [=====] - 27s 274ms/step - loss: 0.6310 - accuracy: 0.6435 - val_loss: 0.6224 - val_accuracy: 0.6560
Epoch 7/20
100/100 [=====] - 29s 291ms/step - loss: 0.5957 - accuracy: 0.6855 - val_loss: 0.6088 - val_accuracy: 0.6610
```

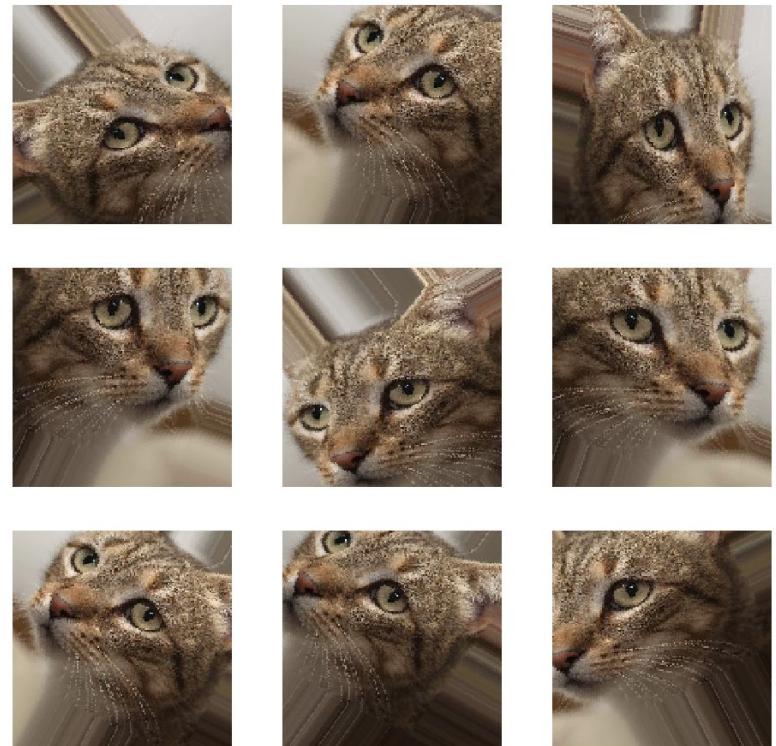
Convolutional Neural Networks - Another Example



Convolutional Neural Networks - Another Example



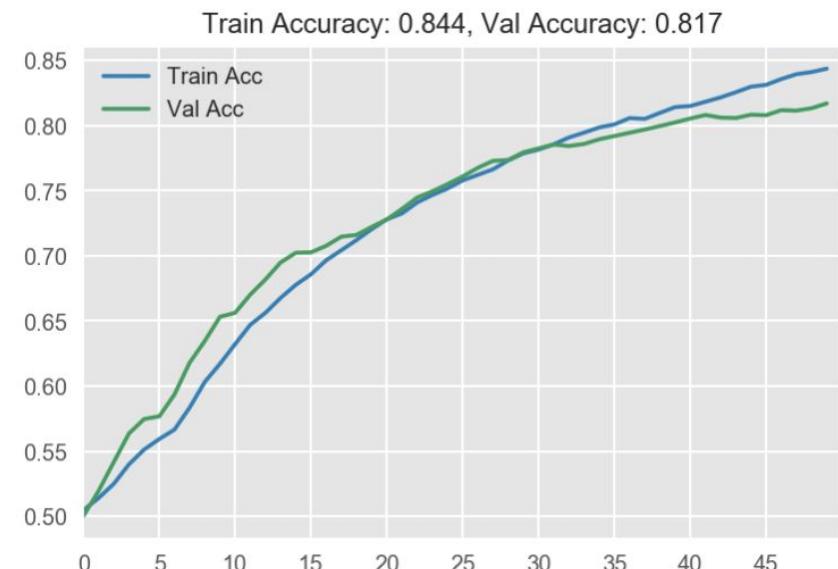
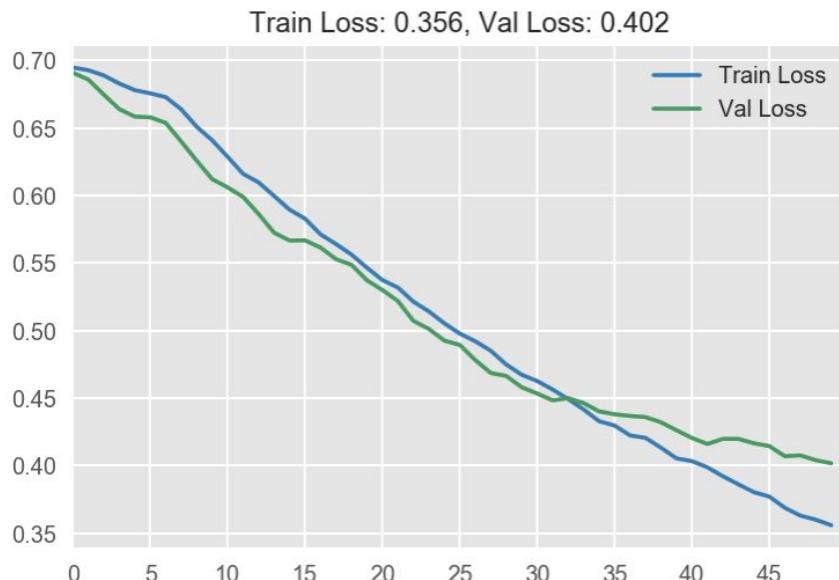
**Data
Augmentation**



Convolutional Neural Networks - Another Example

```
train_datagen = ImageDataGenerator(  
    rescale=1./255,  
    rotation_range=40,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True,)  
  
# Note that the validation data should not be augmented!  
test_datagen = ImageDataGenerator(rescale=1./255)  
  
train_generator = train_datagen.flow_from_directory(  
    train_dir,  
    target_size=(150, 150),  
    batch_size=20,  
    class_mode='binary')  
  
validation_generator = test_datagen.flow_from_directory(  
    validation_dir,  
    target_size=(150, 150),  
    batch_size=20,  
    class_mode='binary')  
  
# early_stop = EarlyStopping(monitor='val_loss', patience=6, verbose=1)  
history_aug = model_aug.fit_generator(train_generator, steps_per_epoch=100, epochs=60,  
                                      validation_data=validation_generator, validation_steps=50, verbose=1)
```

Convolutional Neural Networks - Another Example

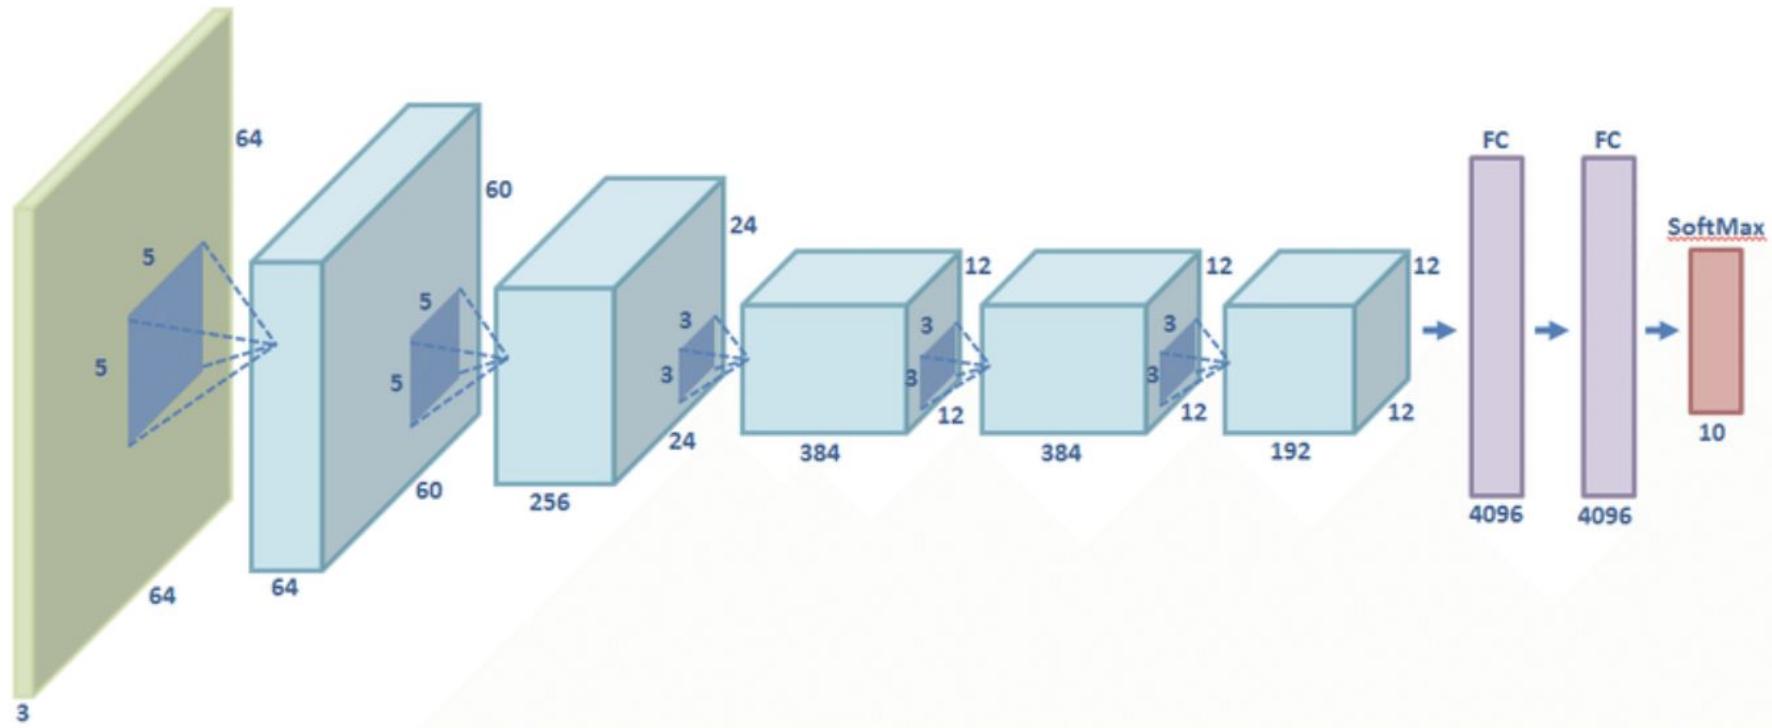


Questions?

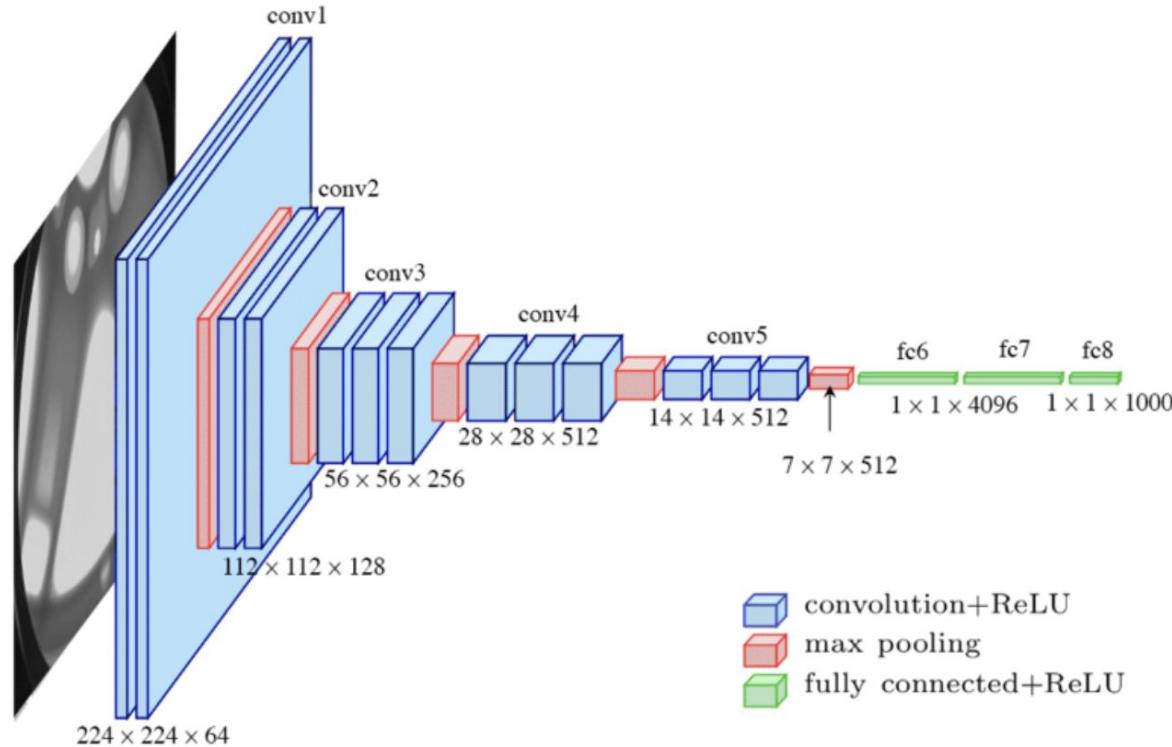
Let's take a 10 min break!

CNN architectures

Convolutional Neural Networks - AlexNet



Convolutional Neural Networks - VGG16



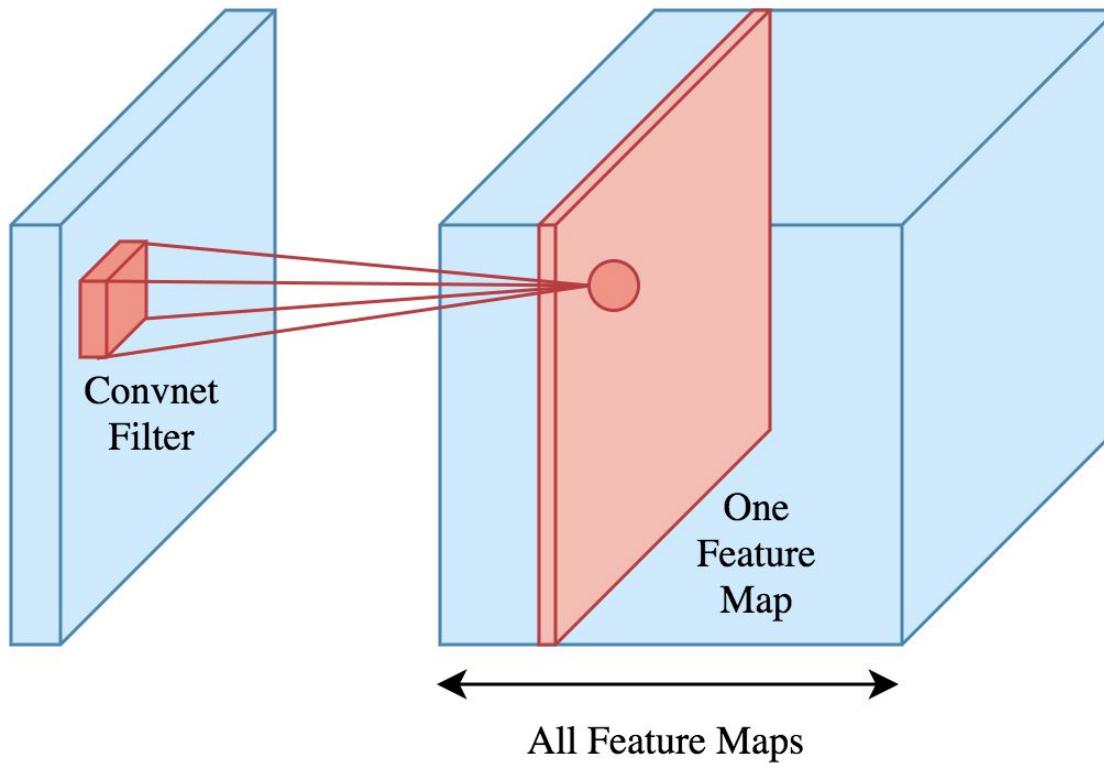
Convolutional Neural Networks - VGG16

- VGG16 uses only 3x3 convolutions
 - Two stacked 3x3 convolutions -> a single 5x5 convolution.
 - Three stacked 3x3 convolutions -> a single 7x7 convolution.
- Stacking two convolutions uses two relu operations, and thus capture more non-linear interactions
- The number of filters increase as we go deeper into the network.
- Trained on 4 GPUs for 3 weeks

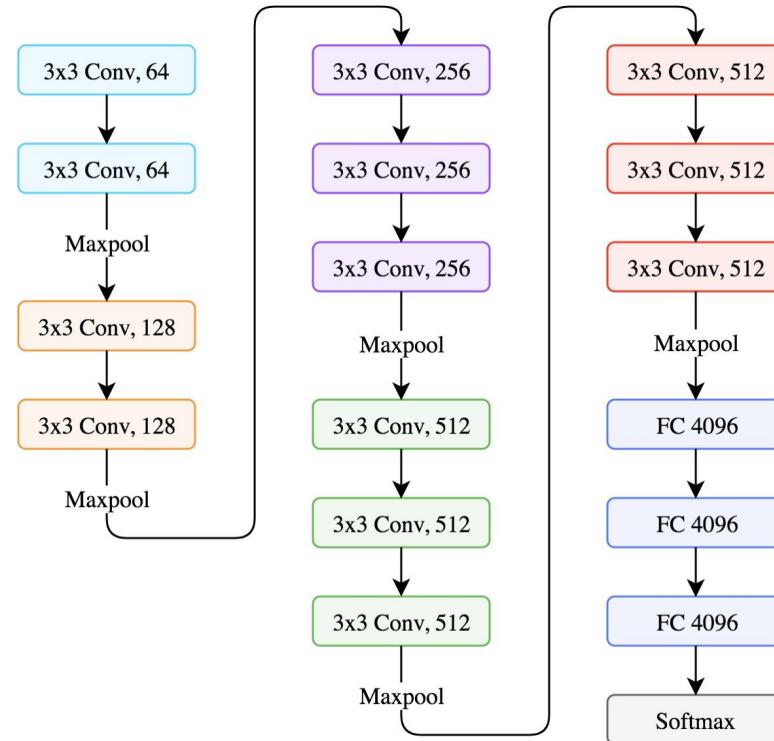
Convolutional Neural Networks - Visualizing VGG16

- Feature maps
- Convnet filters
- Class outputs

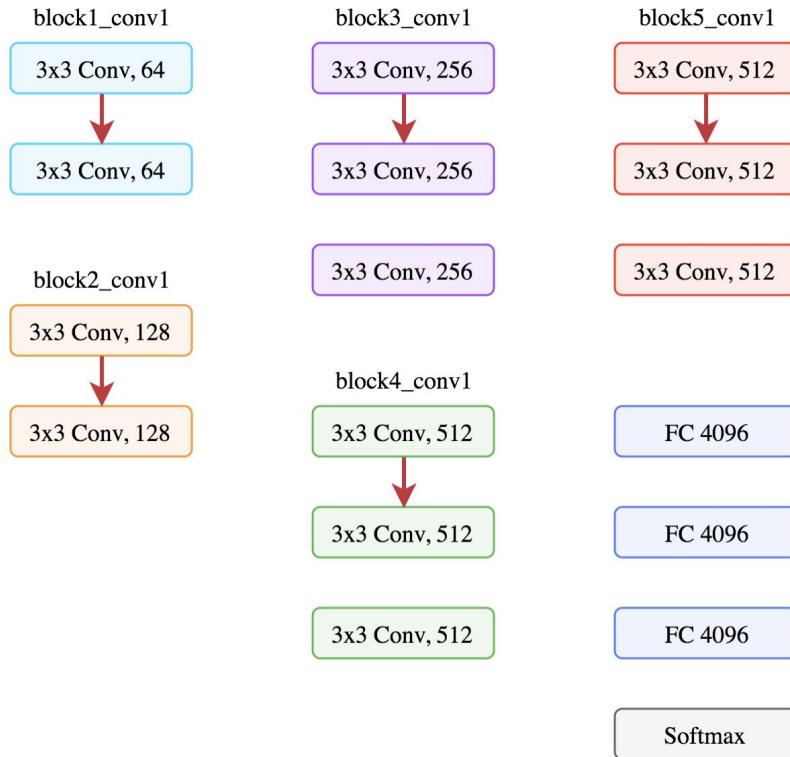
Convolutional Neural Networks - Visualizing VGG16



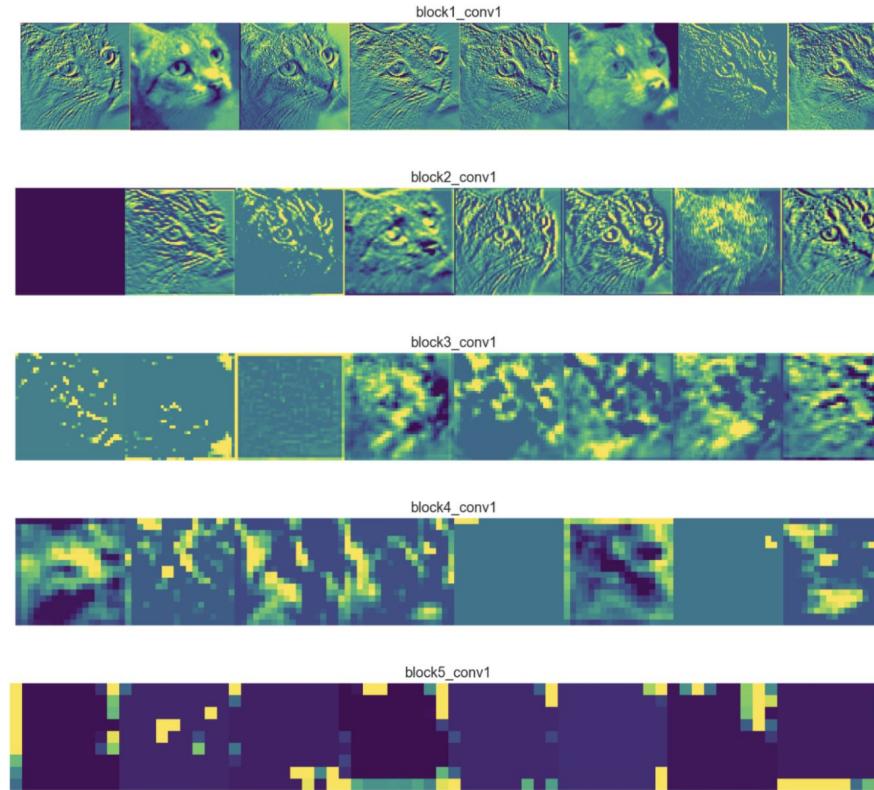
Convolutional Neural Networks - Visualizing VGG16



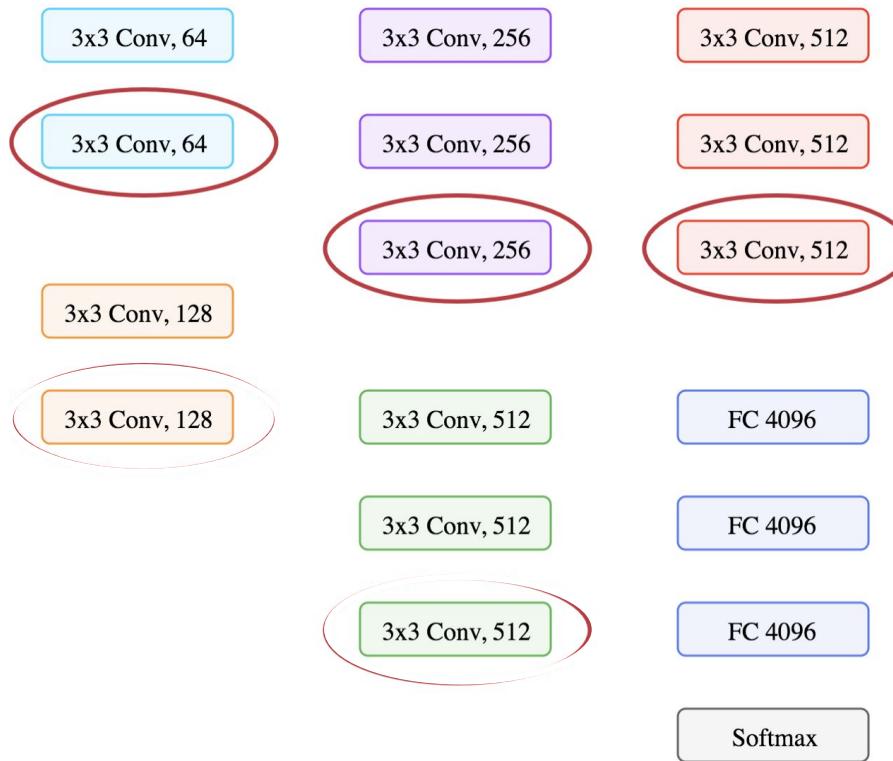
Convolutional Neural Networks - Visualizing Feature Maps



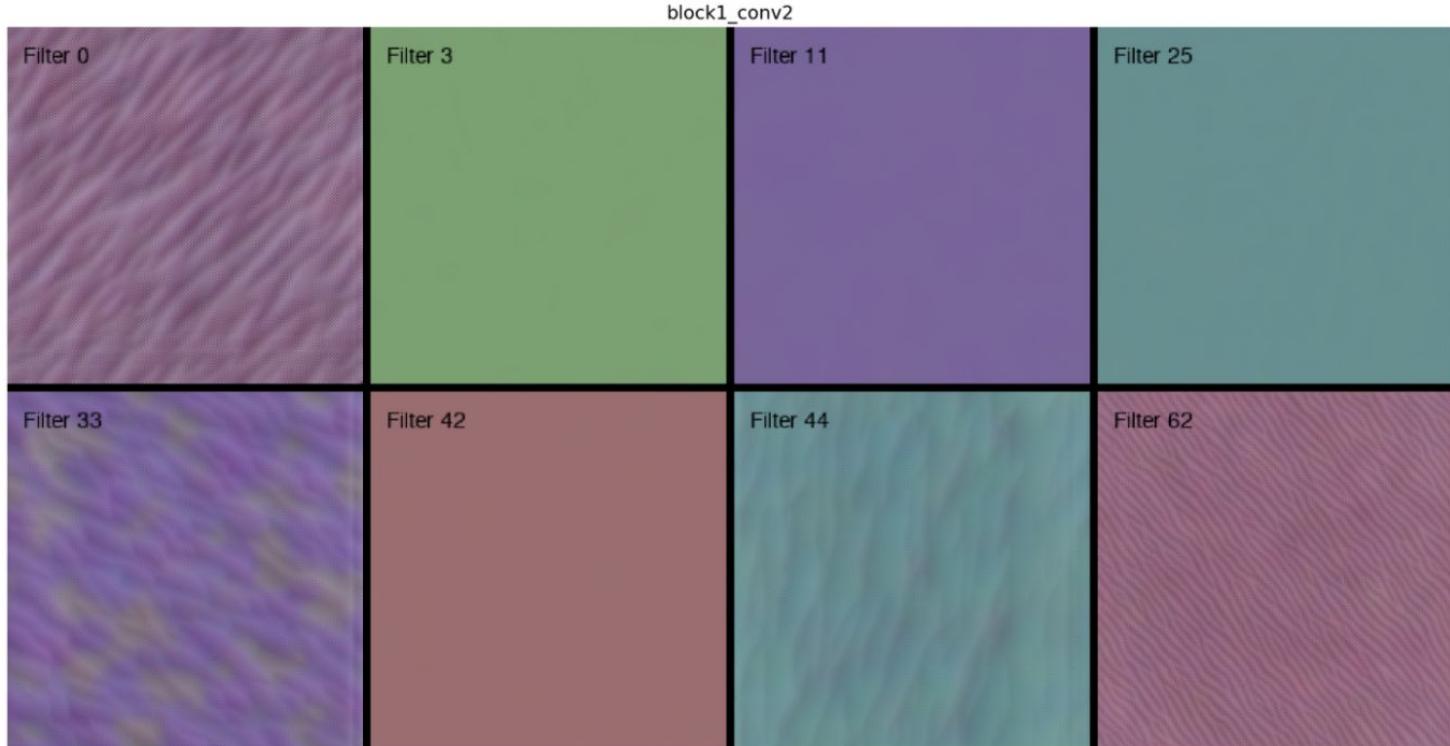
Convolutional Neural Networks - VGG16 Feature Maps



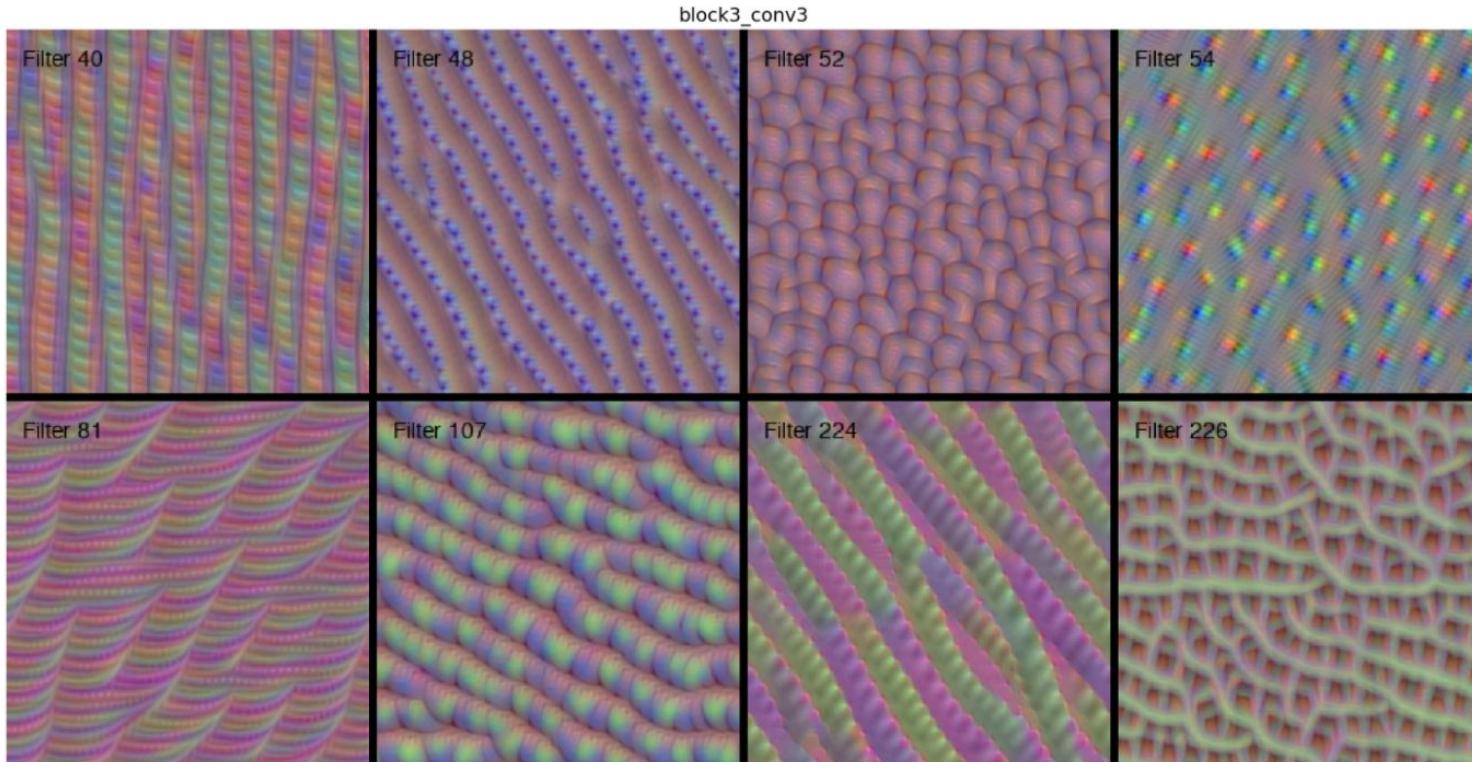
Convolutional Neural Networks - Visualizing Convnet Filters



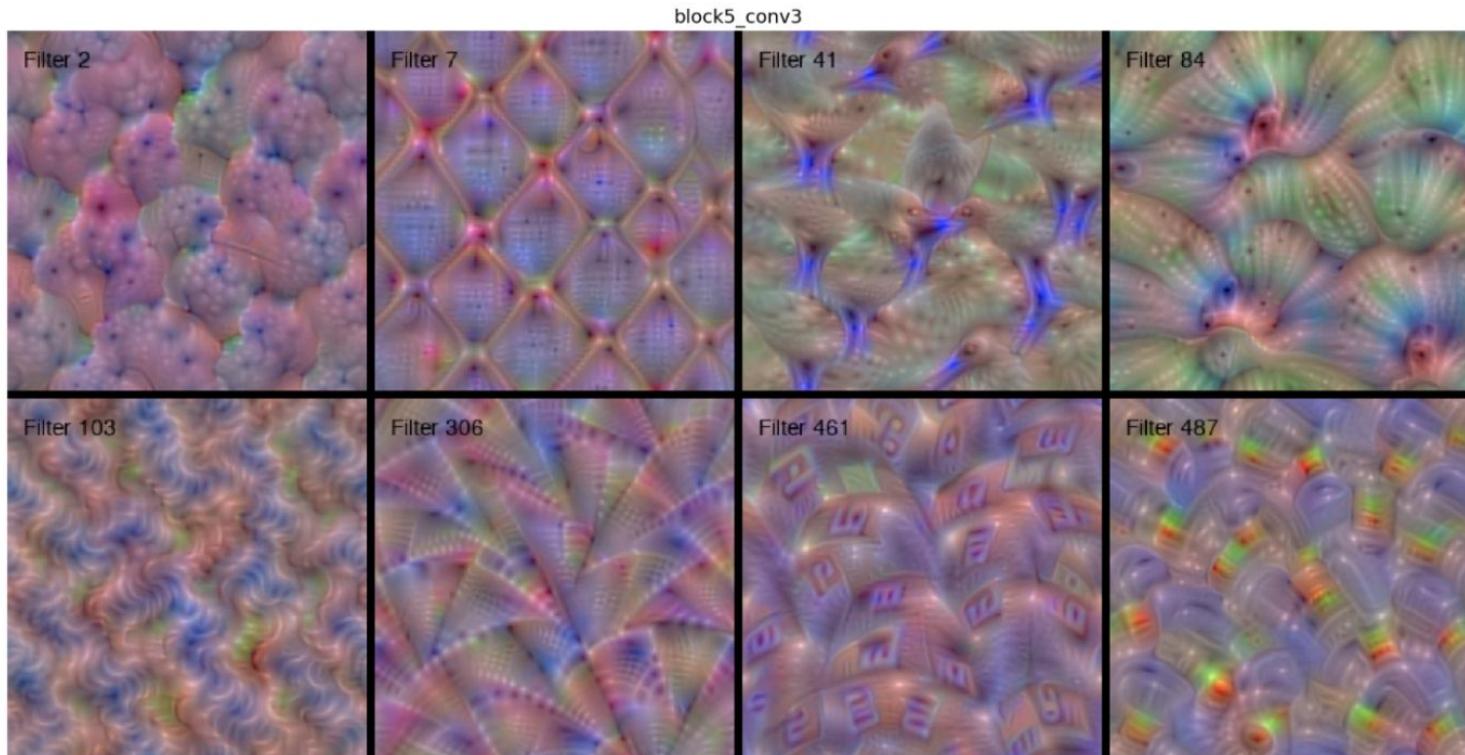
Convolutional Neural Networks - Visualizing Convnet Filters



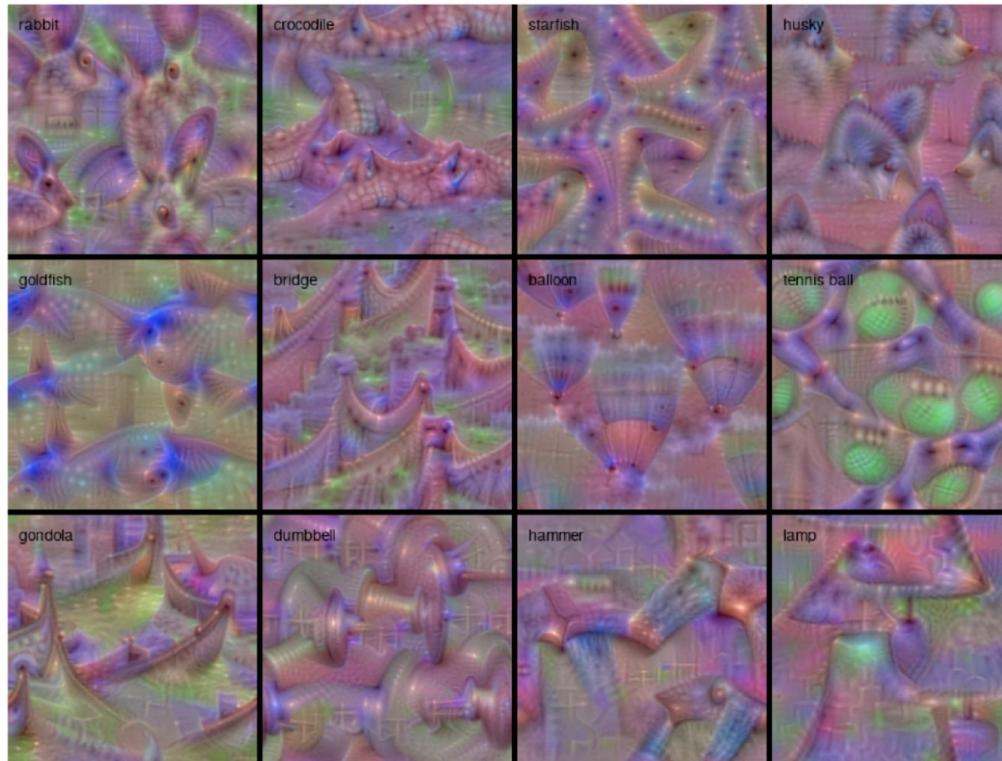
Convolutional Neural Networks - Visualizing Convnet Filters



Convolutional Neural Networks - Visualizing Convnet Filters



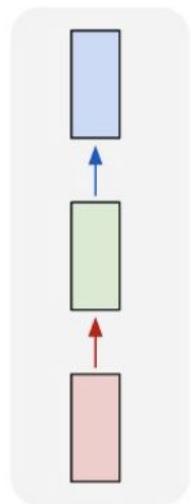
Convolutional Neural Networks - Visualizing Output Layers



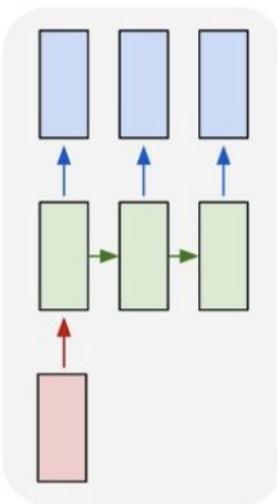
Recurrent Neural Networks

Recurrent Neural Networks

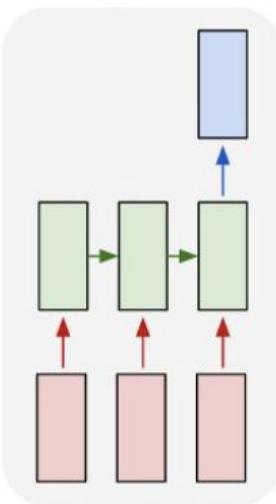
one to one



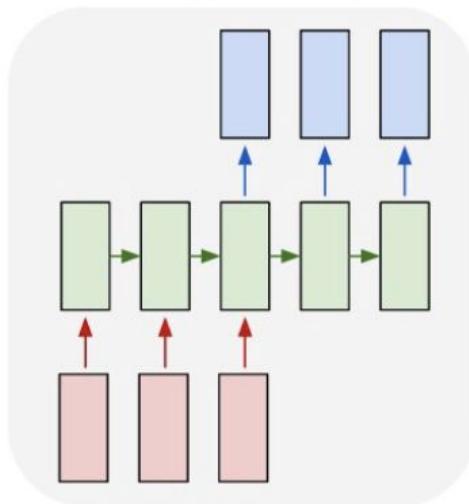
one to many



many to one



many to many



many to many

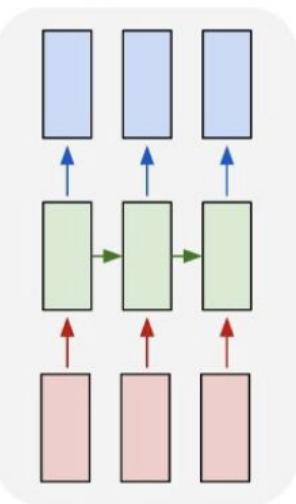


Image
Classification

Image
Captioning

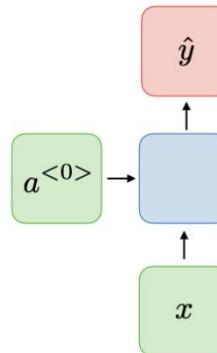
Sentiment
Analysis

Machine
Translation

Video
Classification

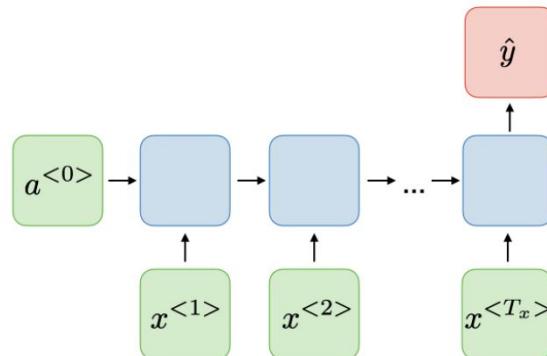
Recurrent Neural Networks

One-to-one
 $T_x = T_y = 1$



Traditional neural network

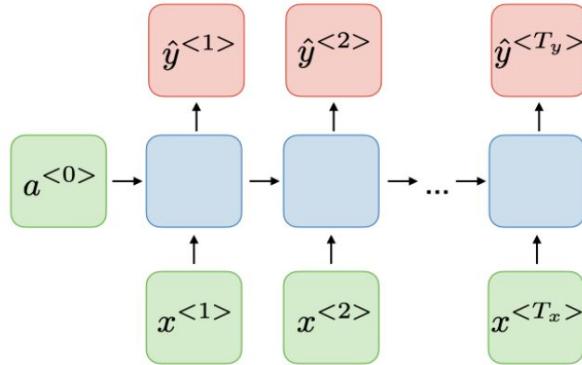
Many-to-one
 $T_x > 1, T_y = 1$



Sentiment classification

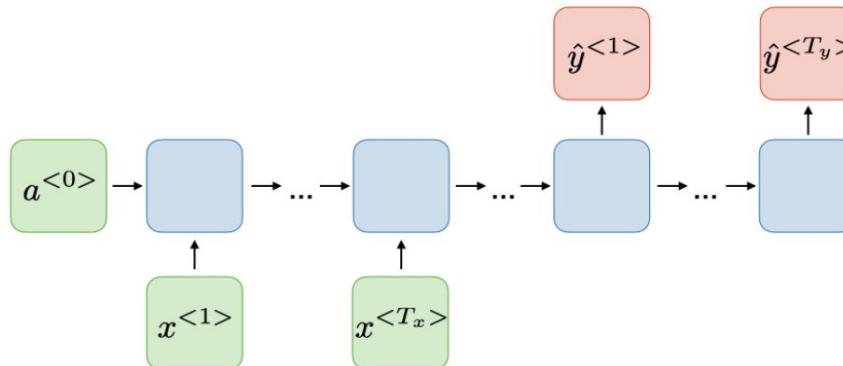
Recurrent Neural Networks

Many-to-many
 $T_x = T_y$



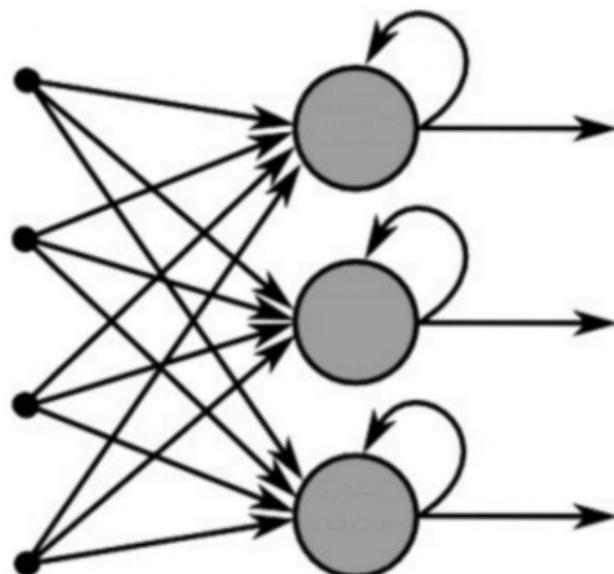
Name entity recognition

Many-to-many
 $T_x \neq T_y$

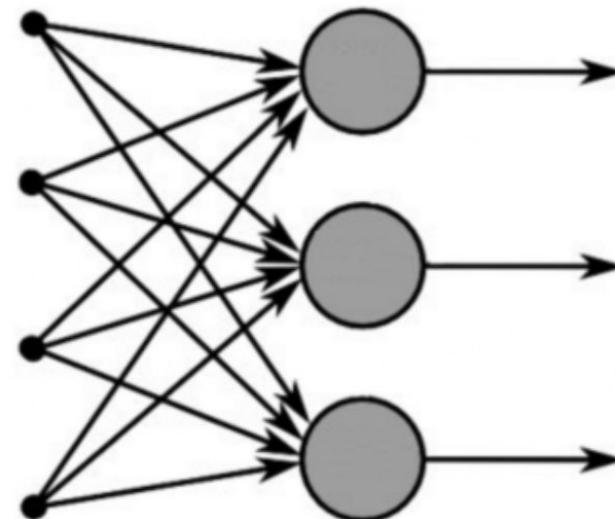


Machine translation

Recurrent Neural Networks

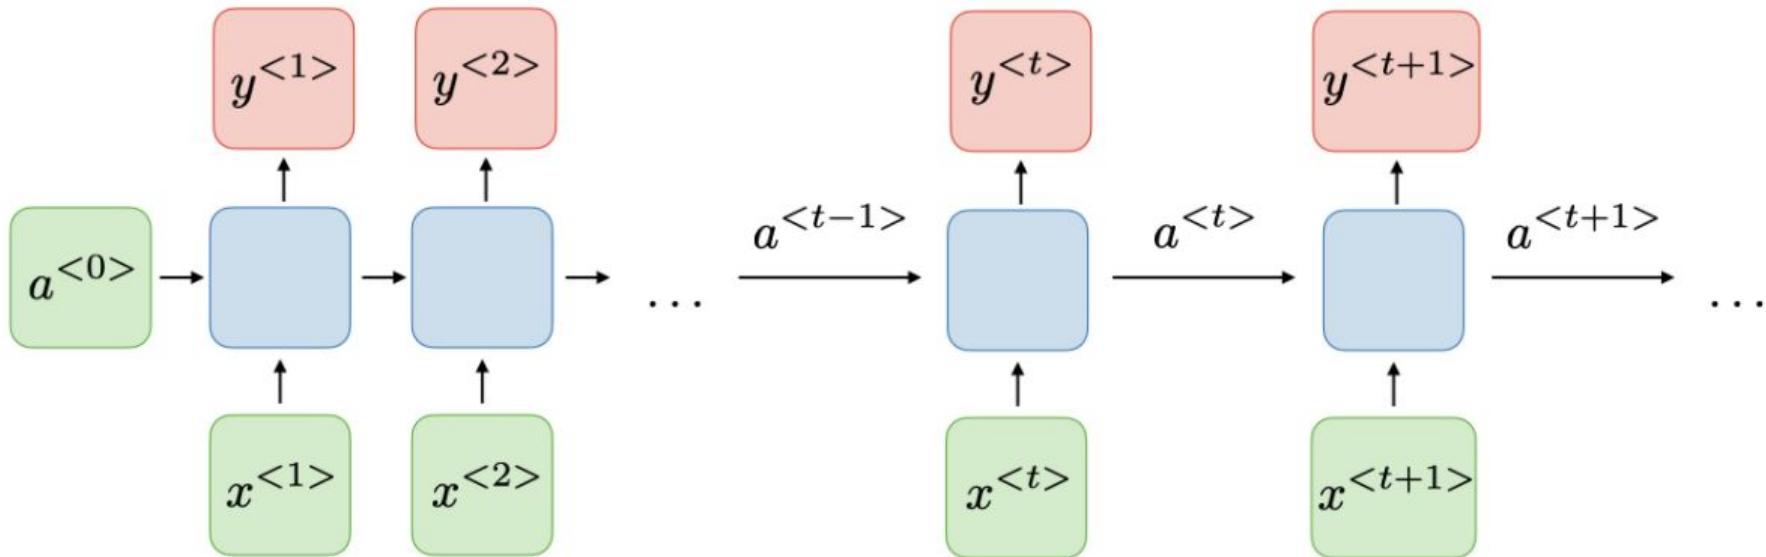


Recurrent Neural Network



Feed-Forward Neural Network

Recurrent Neural Networks

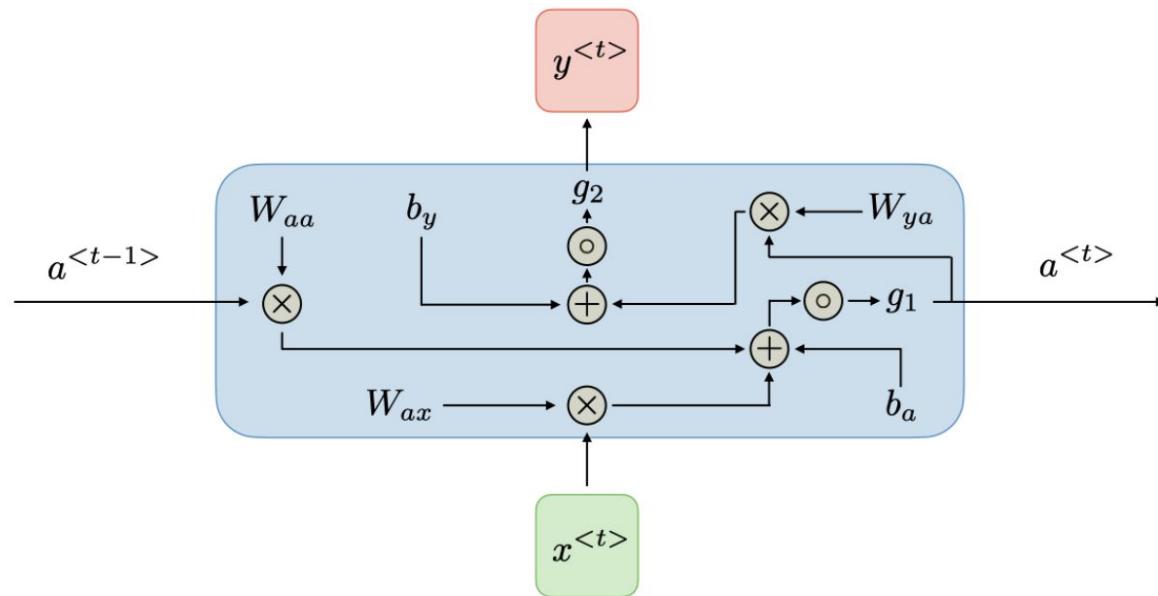


Recurrent Neural Networks

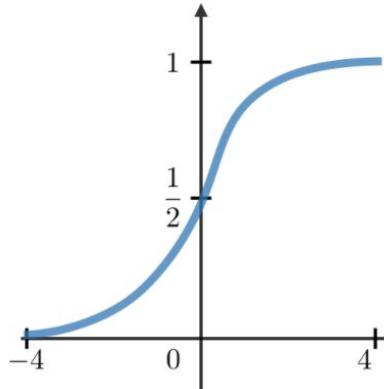
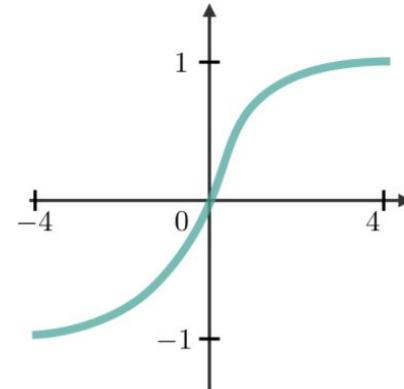
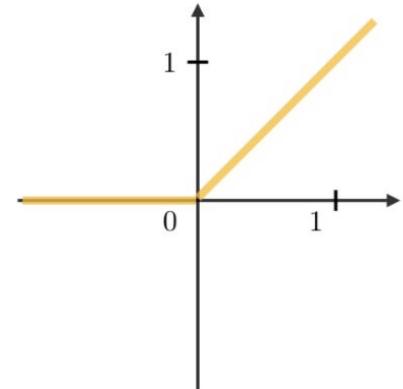
$$a^{<t>} = g_1(W_{aa}a^{<t-1>} + W_{ax}x^{<t>} + b_a)$$

and

$$y^{<t>} = g_2(W_{ya}a^{<t>} + b_y)$$



Recurrent Neural Networks - Activation Functions

| Sigmoid | Tanh | RELU |
|---|---|---|
| $g(z) = \frac{1}{1 + e^{-z}}$  | $g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$  | $g(z) = \max(0, z)$  |

Recurrent Neural Networks - Training

$$\mathcal{L}(\hat{y}, y) = \sum_{t=1}^{T_y} \mathcal{L}(\hat{y}^{}, y^{})$$

Loss function

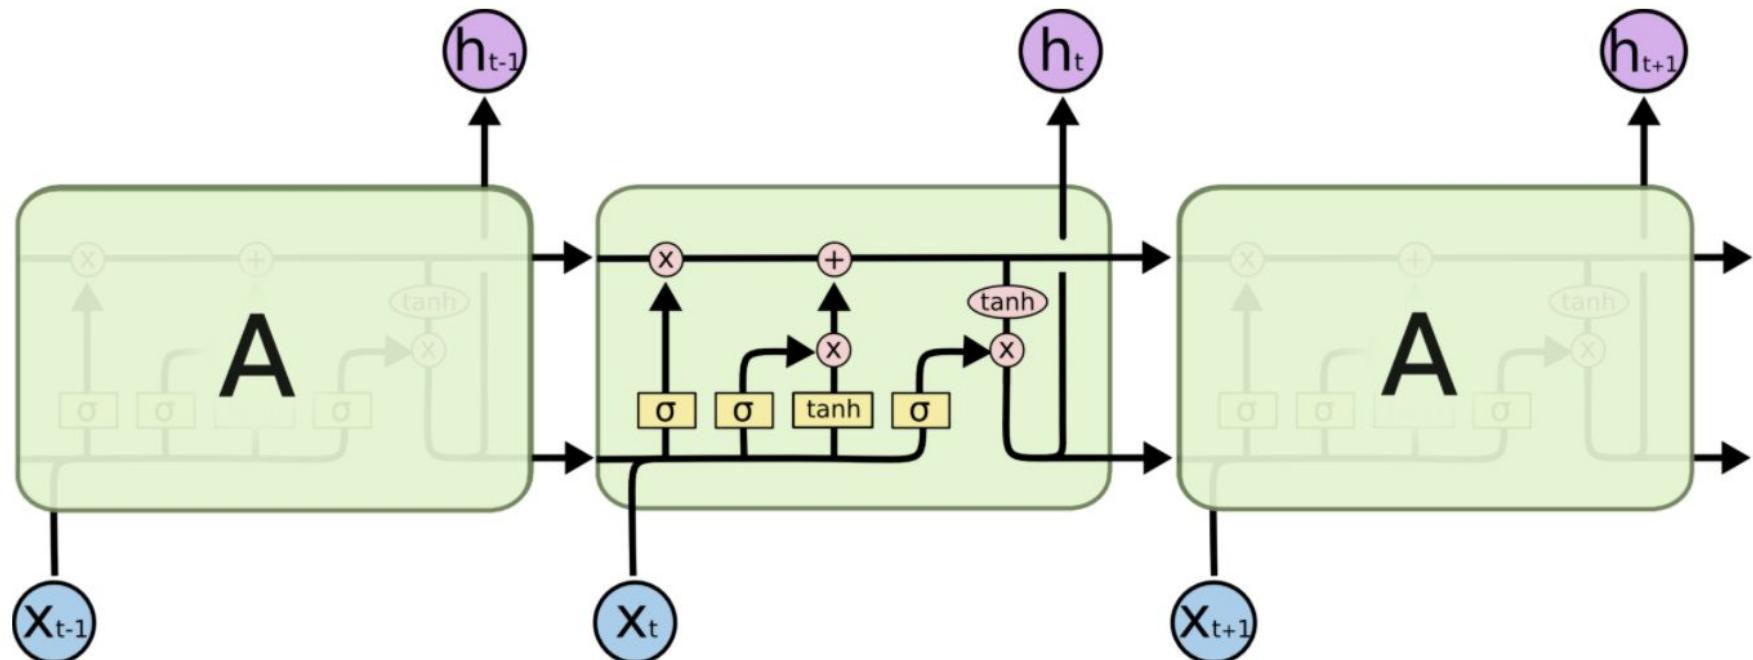
$$\frac{\partial \mathcal{L}^{(T)}}{\partial W} = \sum_{t=1}^T \frac{\partial \mathcal{L}^{(T)}}{\partial W} \Big|_{(t)}$$

Backpropagation through time

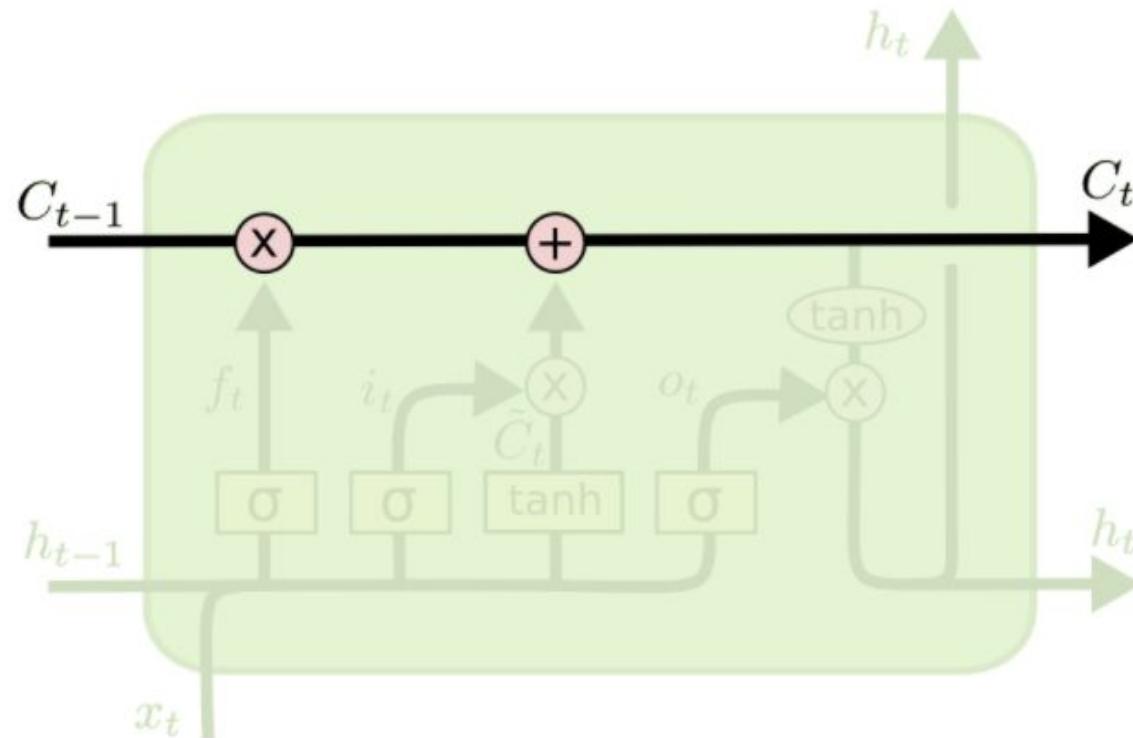
Recurrent Neural Networks - Limitations

- Learning long-term dependencies
 - “I grew up in France... I speak fluent *French*.”
- Vanishing/exploding gradients

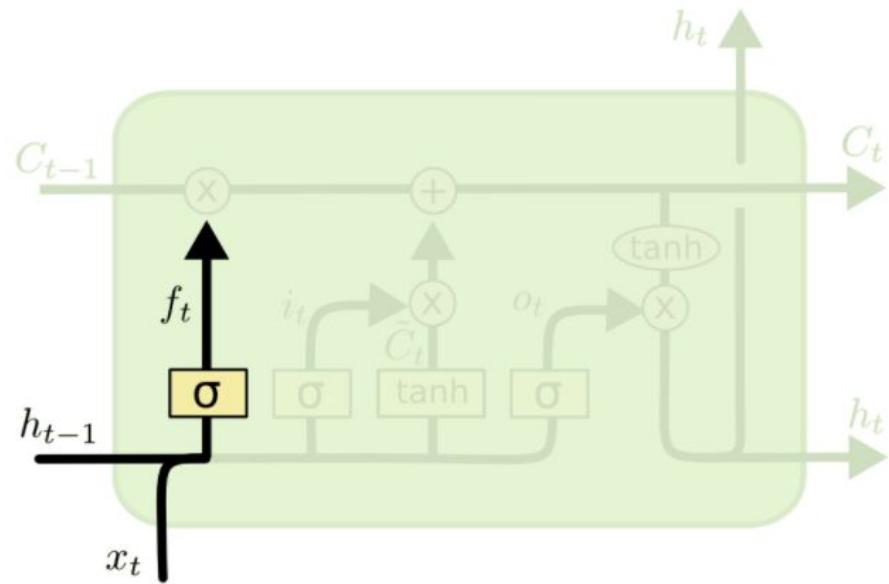
Recurrent Neural Networks - LSTMs



Recurrent Neural Networks - LSTMs



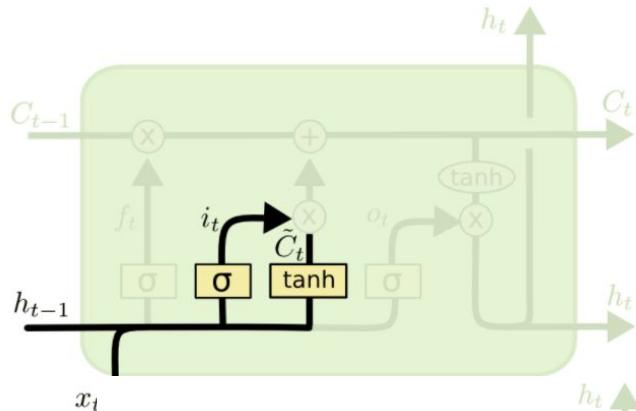
Recurrent Neural Networks - LSTMs



Forget gate

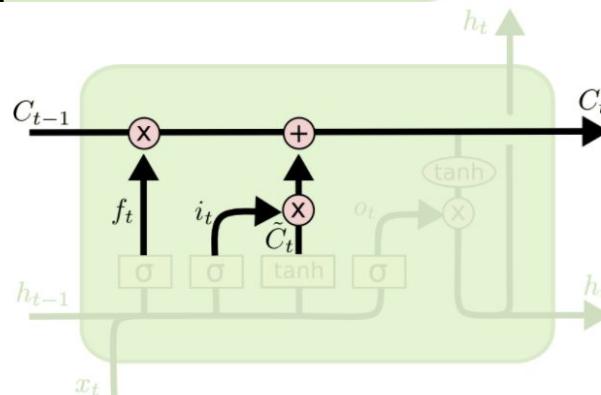
$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$

Recurrent Neural Networks - LSTMs



$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

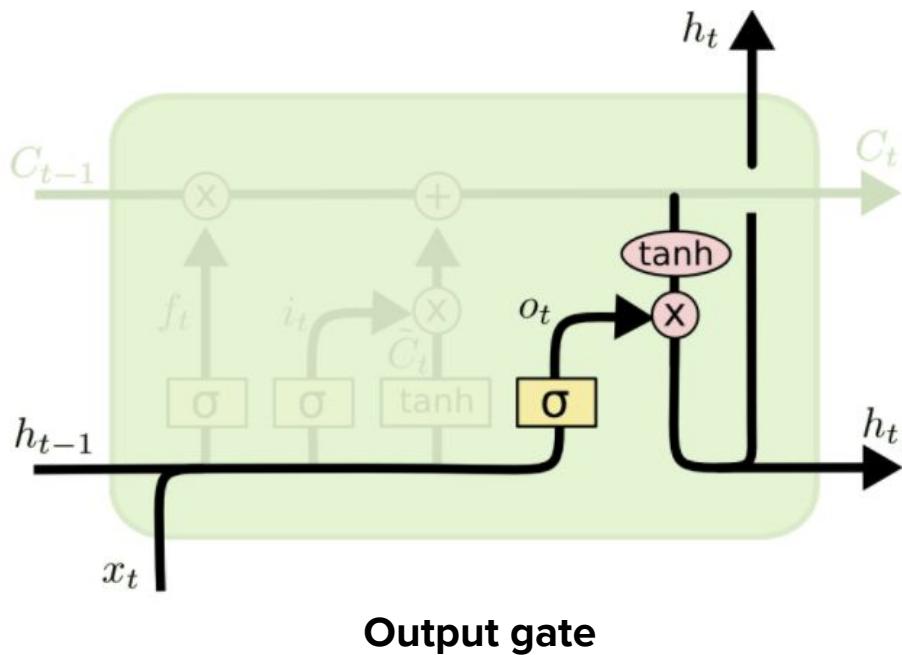
$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

input gate

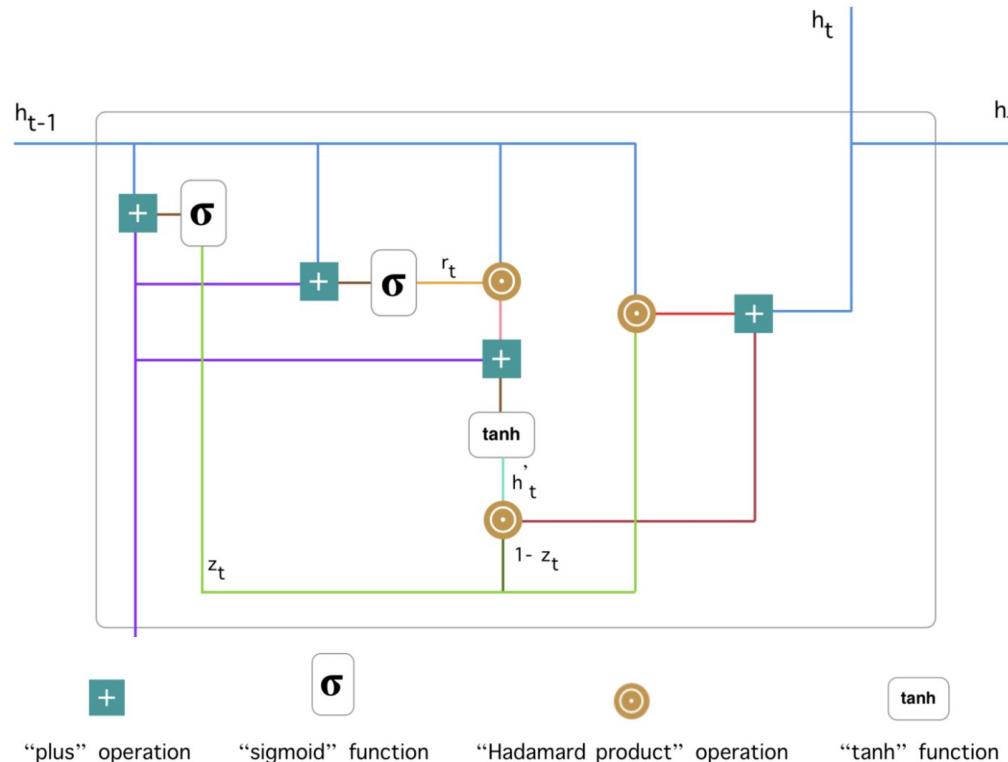
Recurrent Neural Networks - LSTMs



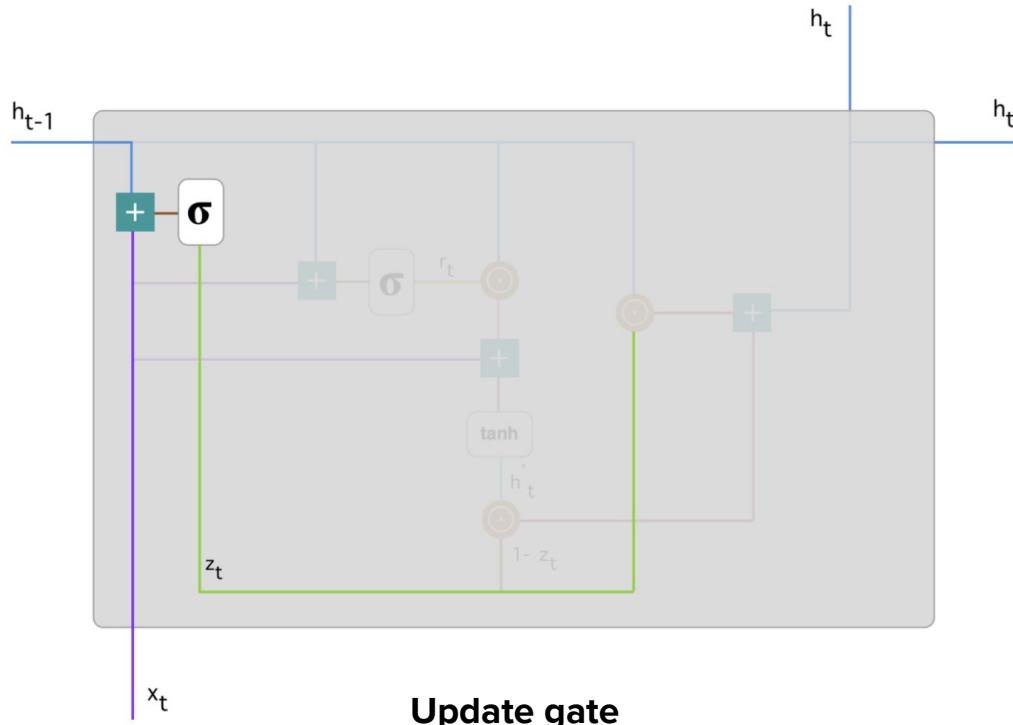
$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh (C_t)$$

Recurrent Neural Networks - GRUs

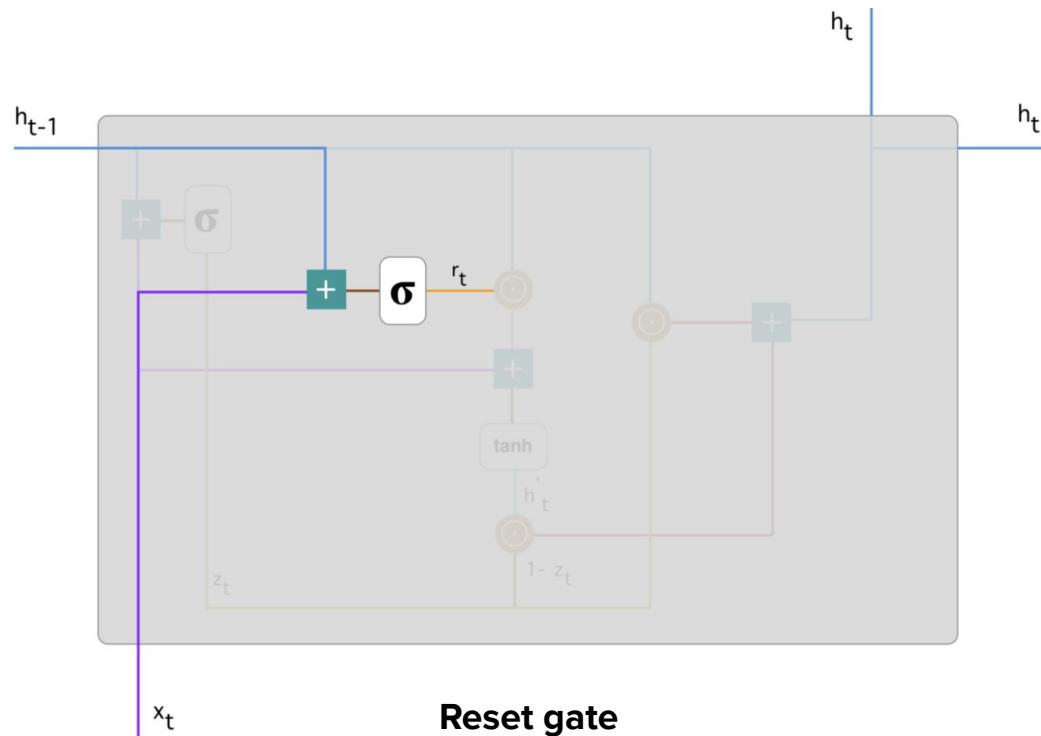


Recurrent Neural Networks - GRUs



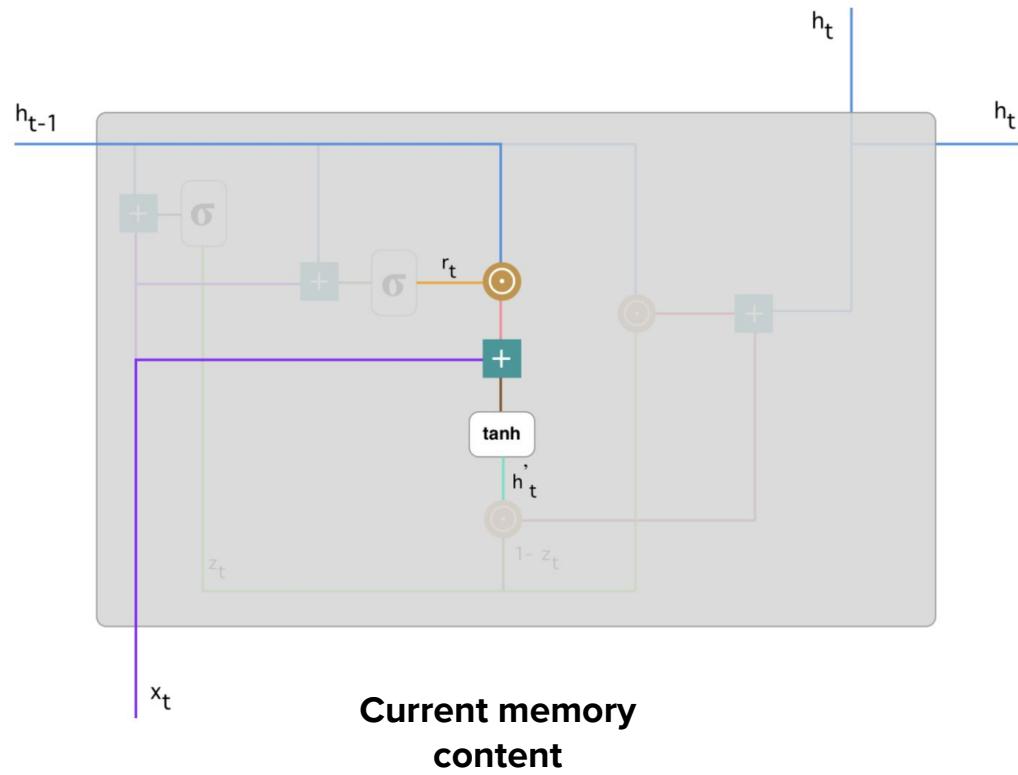
$$z_t = \sigma(W^{(z)}x_t + U^{(z)}h_{t-1})$$

Recurrent Neural Networks - GRUs



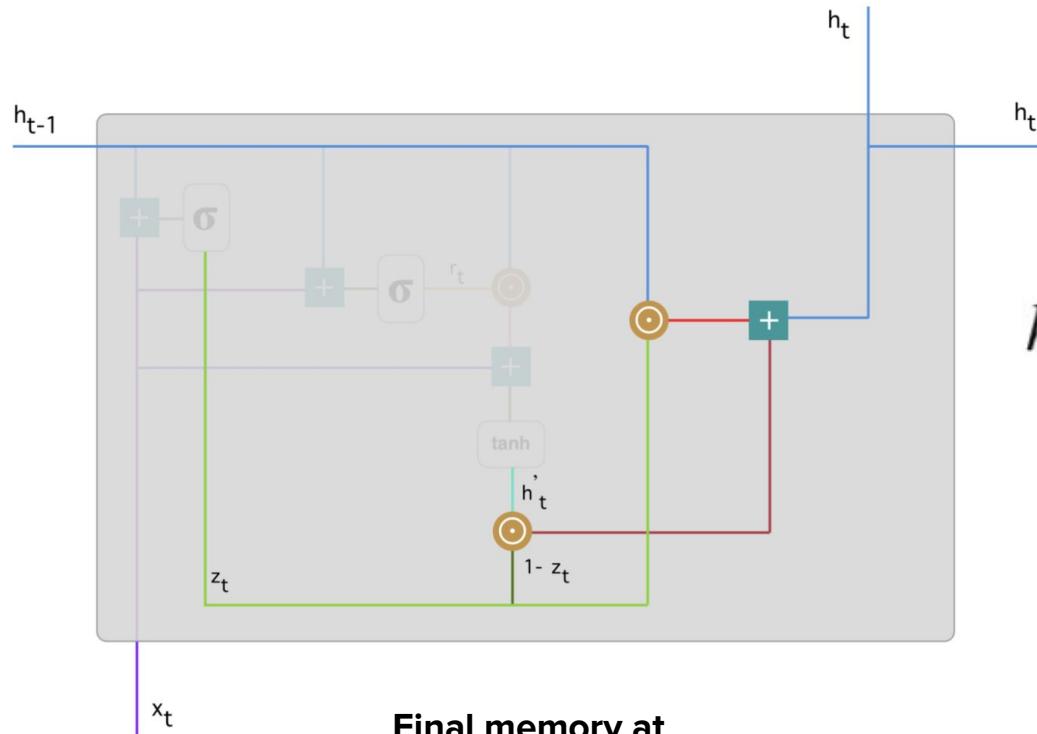
$$r_t = \sigma(W^{(r)}x_t + U^{(r)}h_{t-1})$$

Recurrent Neural Networks - GRUs



$$h'_t = \tanh(Wx_t + r_t \odot Uh_{t-1})$$

Recurrent Neural Networks - GRUs



**Final memory at
current step**

$$h_t = z_t \odot h_{t-1} + (1 - z_t) \odot h'_t$$

Questions?